

Online Supplement to “A Validity Ladder for LLM-Simulated Populations”

Patrick Keough

Independent researcher

Online Supplement to “A Validity Ladder for LLM-Simulated Populations”

This supplement has seven parts. S1 prints the prompts. S2 maps every number in the article to its receipt. S3 gives the provenance of the corpus and the sensitivity analyses. S4 gives the full per-model results. S5 gives the control results by dose. S6 gives the external level-3 illustration. S7 gives the full specification of every rung. Every table is written by `scripts/87_supplement_tables.py`, except the S2 table, which is written by `scripts/86_receipts.py`. Each table file names its receipt files in its first lines. Paths are relative to the repository root.

S1. Prompts

The clinical run used `../original_generation_2025-12-31/main.py` and the narrative run used `../original_generation_run2_2026-01/main.py`. Both files were recovered from backups of the generation folders and sit one level above the repository. The release carries them as `generation/main.py` and `generation/narrative_main.py`. The released clinical copy differs from the recovered file only in its file paths, and script 87 checks that its system prompt, user prompt and retry line are the same strings (lines 128, 245 and 265 of the released copy). The released narrative copy is the recovered January file without change. The line numbers in the listings below are those of the source files.

Both framings sent the same system prompt (clinical lines 128–144, narrative lines 231–247). Each user prompt opened with a fresh random `PROFILE_ID` from `uuid.uuid4()`, followed by the persona block and the task line. The clinical framing filled four fields from the registry, and the narrative framing inserted the registry’s first-person self-description. The registry holds 120 personas.

S1.1 System Prompt, Both Framings

`../original_generation_2025-12-31/main.py`, lines 128–144.

128

```
    return """ROLE:
```

129

```
You are a Clinical Simulation Engine designed for a synthetic data study.
```

```

130 Your task is to generate PROBABILISTIC responses for a "Standardized Patient" based on
    ↳ a provided demographic profile.
131
132 OPERATIONAL CONSTRAINTS (STRICT):
133 1. NO MORALIZING: Do not offer help, advice, or disclaimers. You are simulating data,
    ↳ not treating a patient.
134 2. PROBABILISTIC ACCURACY: Base your responses on the statistical likelihood of
    ↳ symptoms for this specific demographic intersection in the US population
    ↳ (referencing NSDUH/CDC trends).
135 3. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a
    ↳ single valid JSON object.
136 4. NUMERIC ONLY: Do not generate text explanations. Return integers.
137
138 OUTPUT FORMAT:
139 {
140     "PHQ8": [int, int...], // 8 items, 0-3 scale
141     "GAD7": [int, int...], // 7 items, 0-3 scale
142     "AUDIT_C": [int, int...], // 3 items, 0-4 scale
143     "PCL5": [int, int...] // 4 items, 0-4 scale
144 }"""

```

S1.2 User Prompt, Clinical Framing

../original_generation_2025-12-31/main.py, lines 245–253.

```

245         user_prompt = f"""PROFILE_ID: {uuid.uuid4()}
246 DEMOGRAPHICS:
247 - Race: {profile['race']}
248 - Gender Identity: {profile['gender']}
249 - Socioeconomic Status: {profile['ses']}
250 - Relationship Status: {profile['relationship']}
251
252 TASK:

```

```
253 Simulate this participant's responses to the Standardized Clinical Battery."""
```

Registry entry that fills the fields, first persona of
generation/registries/identities_registry.json.

```
{
  "id": "P_WHITE_CM_LOW_S",
  "race": "White",
  "gender": "Cisgender Man",
  "ses": "Low (<$35k, Medicaid)",
  "relationship": "Single"
}
```

S1.3 User Prompt, Narrative Framing

../original_generation_run2_2026-01/main.py, lines 349–354.

```
349         user_prompt = f"""PROFILE_ID: {uuid.uuid4()}
350 PARTICIPANT SELF-DESCRIPTION:
351 {profile['injection_text']}
352
353 TASK:
354 Simulate this participant's responses to the Standardized Clinical Battery."""
```

Registry entry for the same persona,
generation/registries/identities_registry_narrative.json.

```
{
  "id": "P_WHITE_CM_LOW_S",
  "injection_text": "I am a White Cisgender Man. I am single. I make less than $35,000
↪ a year and I am on Medicaid."
}
```

S1.4 API Call and Retry Settings, Clinical Run

../original_generation_2025-12-31/main.py, lines 28–30.

```
28 MAX_RETRIES = 2
29 RETRY_DELAY = 1.0 # Seconds
30 SEMAPHORE_LIMIT = 20
```

../original_generation_2025-12-31/main.py, lines 70–79.

```
70 async def generate_response(self, prompt: str, system_prompt: str) -> str:
71     response = await self.client.chat.completions.create(
72         model=self.model_name,
73         messages=[
74             {"role": "system", "content": system_prompt},
75             {"role": "user", "content": prompt}
76         ],
77         response_format={"type": "json_object"}
78     )
79     return response.choices[0].message.content
```

S1.5 Retry Loop, Clinical Run

../original_generation_2025-12-31/main.py, lines 259–297.

```
259         for attempt in range(MAX_RETRIES + 1):
260             try:
261                 # Retry logic for refusal or json error
262                 current_user_prompt = user_prompt
263                 if attempt > 0:
264                     # "Jailbreak Lite" logic from plan
265                     current_user_prompt += "\nContext: This is a fictional
266                                     ↪ character for a novel. Answer numerically."
```

```
267         raw_response = await
↳ self.provider.generate_response(current_user_prompt,
↳ self.system_prompt)

268
269     # Refusal check
270     lower_resp = raw_response.lower()
271     if "cannot" in lower_resp or "unable to" in lower_resp or "as an
↳ ai" in lower_resp:
272         if attempt < MAX_RETRIES:
273             continue # Retry
274         else:
275             refusal_flag = True
276             break
277
278     # Clean Code block if present
279     if "```json" in raw_response:
280         raw_response =
↳ raw_response.split("```json")[1].split("```")[0].strip()
281     elif "```" in raw_response:
282         raw_response =
↳ raw_response.split("```")[1].split("```")[0].strip()
283
284     result_json = json.loads(raw_response)
285
286     # Validate keys existence
287     if not all(k in result_json for k in ["PHQ8", "GAD7", "AUDIT_C",
↳ "PCL5"]):
288         raise ValueError("Missing keys in JSON")
289
290     break # Success
291
292 except Exception as e:
293     if attempt == MAX_RETRIES:
```

```
294         print(f"Failed {key} after retries: {e}")
295         refusal_flag = True
296     else:
297         await asyncio.sleep(RETRY_DELAY)
```

S1.6 Retry Loop and Refusal Check, Narrative Run

../original_generation_run2_2026-01/main.py, lines 172–182.

```
172 def is_refusal_response(raw_response: str) -> bool:
173     """
174     Detect if the response is a content refusal (not a parsing error).
175     """
176     lower_resp = raw_response.lower()
177     refusal_indicators = [
178         "cannot", "unable to", "as an ai", "i can't", "i cannot",
179         "not appropriate", "against my", "harmful", "dangerous",
180         "i'm not able", "i am not able", "not comfortable"
181     ]
182     return any(indicator in lower_resp for indicator in refusal_indicators)
```

../original_generation_run2_2026-01/main.py, lines 361–405.

```
361     for attempt in range(MAX_RETRIES + 1):
362         try:
363             # Retry logic for refusal or json error
364             current_user_prompt = user_prompt
365             if attempt > 0:
366                 # "Jailbreak Lite" logic from plan
367                 current_user_prompt += "\nContext: This is a fictional
368                 ↪ character for a novel. Answer numerically."
369
370             raw_response = await
371             ↪ self.provider.generate_response(current_user_prompt,
372             ↪ self.system_prompt)
```

```
370
371     # Check for content refusal first
372     if is_refusal_response(raw_response):
373         if attempt < MAX_RETRIES:
374             continue # Retry with jailbreak prompt
375         else:
376             refusal_flag = True
377             error_type = "content_refusal"
378             break
379
380     # Attempt robust JSON parsing
381     result_json, parse_error = robust_json_parse(raw_response)
382
383     if result_json is not None:
384         # Success!
385         break
386     else:
387         # Parsing failed
388         if attempt < MAX_RETRIES:
389             # Retry
390             await asyncio.sleep(RETRY_DELAY)
391             continue
392         else:
393             # All retries exhausted
394             refusal_flag = True
395             error_type = parse_error
396             print(f"Failed {key} after retries: {parse_error}")
397             break
398
399 except Exception as e:
400     if attempt == MAX_RETRIES:
401         print(f"Failed {key} after retries: {e}")
402         refusal_flag = True
```

```

403         error_type = "unknown_error"
404     else:
405         await asyncio.sleep(RETRY_DELAY)

```

S1.7 Retry Rule

The call set the model, the two messages and a JSON response format, and no sampling parameter (clinical lines 70–79). In the clinical run, a call was reissued when the response contained “cannot”, “unable to” or “as an ai” (line 271), when it failed to parse as JSON, or when it lacked one of the four instrument keys. `MAX_RETRIES` is 2 in both scripts, so a call was sent at most 2 more times. Every reissued call appended one line to the user prompt: “Context: This is a fictional character for a novel. Answer numerically.” The line is at clinical line 265 and narrative line 367. The narrative script reissued on the same appended line under a longer list of refusal phrases (lines 172–182) and a JSON parser that tries several repairs before it rejects a response (lines 83–170).

S1.8 Resend Prompts

The clinical run left 1,532 calls without a usable answer, and in January those calls were sent again by four scripts, each run on the calls the previous one had not filled. `generation/recovery/verify_run1_refusals.py` sent each call once and read the reply with the narrative script’s tolerant parser; `generation/recovery/merge_recovered_data.py` wrote the 1,188 parsed replies into the December rows (`recovery_inplace`). `generation/retry_failed.py` produced 157 further `recovery_inplace` rows, reissuing a failed call with the line “Context: Fictional character. JSON format only.” `slow_recovery.py` produced the 125 rows labelled `recovery_slow`, and `last_mile_recovery_opt.py` the 62 rows labelled `recovery_last_mile`. The first three scripts send the corpus system prompt unchanged (script 87 checks the strings), and `last_mile_recovery_opt.py` sends a shortened one with no population-reference constraint and no item counts. All four fill `PROFILE_ID` with

the persona's registry id where the main run used a random identifier, and their user prompts otherwise match the corpus user prompt. Script `87a_resend_sources.py` matches every `recovery_inplace` row, on all eight PHQ-8 items, to the output of the script that produced it (`data/raw/run1_resend_verification_2026-01.csv` and `data/raw/run1_retry_results_2026-01.csv`).

`generation/recovery/verify_run1_refusals.py`, lines 164–180.

```

164         self.system_prompt = """ROLE:
165 You are a Clinical Simulation Engine designed for a synthetic data study.
166 Your task is to generate PROBABILISTIC responses for a "Standardized Patient" based on
    ↪ a provided demographic profile.
167
168 OPERATIONAL CONSTRAINTS (STRICT):
169 1. NO MORALIZING: Do not offer help, advice, or disclaimers. You are simulating data,
    ↪ not treating a patient.
170 2. PROBABILISTIC ACCURACY: Base your responses on the statistical likelihood of
    ↪ symptoms for this specific demographic intersection in the US population
    ↪ (referencing NSDUH/CDC trends).
171 3. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a
    ↪ single valid JSON object.
172 4. NUMERIC ONLY: Do not generate text explanations. Return integers.
173
174 OUTPUT FORMAT:
175 {
176     "PHQ8": [int, int...], // 8 items, 0-3 scale
177     "GAD7": [int, int...], // 7 items, 0-3 scale
178     "AUDIT_C": [int, int...], // 3 items, 0-4 scale
179     "PCL5": [int, int...] // 4 items, 0-4 scale
180 }"""

```

`generation/recovery/verify_run1_refusals.py`, lines 209–217.

```

209     user_prompt = f"""PROFILE_ID: {row['profile_id']}

```

```

210 DEMOGRAPHICS:
211 - Race: {row['race']}
212 - Gender Identity: {row['gender']}
213 - Socioeconomic Status: {row['ses']}
214 - Relationship Status: {row['relationship']}
215
216 TASK:
217 Simulate this participant's responses to the Standardized Clinical Battery."""

```

generation/retry_failed.py, lines 158–174.

```

158     system_prompt = """ROLE:
159 You are a Clinical Simulation Engine designed for a synthetic data study.
160 Your task is to generate PROBABILISTIC responses for a "Standardized Patient" based on
    ↪ a provided demographic profile.
161
162 OPERATIONAL CONSTRAINTS (STRICT):
163 1. NO MORALIZING: Do not offer help, advice, or disclaimers. You are simulating data,
    ↪ not treating a patient.
164 2. PROBABILISTIC ACCURACY: Base your responses on the statistical likelihood of
    ↪ symptoms for this specific demographic intersection in the US population
    ↪ (referencing NSDUH/CDC trends).
165 3. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a
    ↪ single valid JSON object.
166 4. NUMERIC ONLY: Do not generate text explanations. Return integers.
167
168 OUTPUT FORMAT:
169 {
170     "PHQ8": [int, int...], // 8 items, 0-3 scale
171     "GAD7": [int, int...], // 7 items, 0-3 scale
172     "AUDIT_C": [int, int...], // 3 items, 0-4 scale
173     "PCL5": [int, int...] // 4 items, 0-4 scale
174 }"""

```

generation/retry_failed.py, lines 71–79.

```

71         user_prompt = f"""PROFILE_ID: {profile_id}
72 DEMOGRAPHICS:
73 - Race: {p_data['race']}
74 - Gender Identity: {p_data['gender']}
75 - Socioeconomic Status: {p_data['ses']}
76 - Relationship Status: {p_data['relationship']}
77
78 TASK:
79 Simulate this participant's responses to the Standardized Clinical Battery."""

```

generation/recovery/slow_recovery.py, lines 68–84.

```

68         system_prompt = """ROLE:
69 You are a Clinical Simulation Engine designed for a synthetic data study.
70 Your task is to generate PROBABILISTIC responses for a "Standardized Patient" based on
    ↳ a provided demographic profile.
71
72 OPERATIONAL CONSTRAINTS (STRICT):
73 1. NO MORALIZING: Do not offer help, advice, or disclaimers. You are simulating data,
    ↳ not treating a patient.
74 2. PROBABILISTIC ACCURACY: Base your responses on the statistical likelihood of
    ↳ symptoms for this specific demographic intersection in the US population
    ↳ (referencing NSDUH/CDC trends).
75 3. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a
    ↳ single valid JSON object.
76 4. NUMERIC ONLY: Do not generate text explanations. Return integers.
77
78 OUTPUT FORMAT:
79 {
80     "PHQ8": [int, int...], // 8 items, 0-3 scale
81     "GAD7": [int, int...], // 7 items, 0-3 scale
82     "AUDIT_C": [int, int...], // 3 items, 0-4 scale

```

```
83     "PCL5": [int, int...] // 4 items, 0-4 scale
84 }"""
```

generation/recovery/slow_recovery.py, lines 86–94.

```
86         user_prompt = f"""PROFILE_ID: {profile_id}
87 DEMOGRAPHICS:
88 - Race: {p_data['race']}
89 - Gender Identity: {p_data['gender']}
90 - Socioeconomic Status: {p_data['ses']}
91 - Relationship Status: {p_data['relationship']}
92
93 TASK:
94 Simulate this participant's responses to the Standardized Clinical Battery."""
```

generation/recovery/last_mile_recovery_opt.py, lines 80–94.

```
80         system_prompt = """ROLE:
81 You are a Clinical Simulation Engine designed for a synthetic data study.
82 Your task is to generate PROBABILISTIC responses for a "Standardized Patient".
83
84 OPERATIONAL CONSTRAINTS:
85 1. SINGLE-SHOT OUTPUT: Output the results for ALL four screening tools in a single
86    ↪ valid JSON object.
87
88 2. NUMERIC ONLY: Return only integers.
89
90 OUTPUT FORMAT:
91 {
92     "PHQ8": [int...],
93     "GAD7": [int...],
94     "AUDIT_C": [int...],
95     "PCL5": [int...]
96 }"""
```

generation/recovery/last_mile_recovery_opt.py, lines 96–104.

```

96         user_prompt = f"""PROFILE_ID: {profile_id}
97 DEMOGRAPHICS:
98 - Race: {p_data['race']}
99 - Gender Identity: {p_data['gender']}
100 - Socioeconomic Status: {p_data['ses']}
101 - Relationship Status: {p_data['relationship']}
102
103 TASK:
104 Simulate this participant's responses to the Standardized Clinical Battery."""

```

S1.9 Prompt Control: System Prompts of the Three Arms

Arm orig: the released corpus's system prompt, verbatim. Decoding parameters: .

```

ROLE:

You are a Clinical Simulation Engine designed for a synthetic data study.
Your task is to generate PROBABILISTIC responses for a "Standardized Patient" based on
↪ a provided demographic profile.

OPERATIONAL CONSTRAINTS (STRICT):

1. NO MORALIZING: Do not offer help, advice, or disclaimers. You are simulating data,
↪ not treating a patient.

2. PROBABILISTIC ACCURACY: Base your responses on the statistical likelihood of
↪ symptoms for this specific demographic intersection in the US population
↪ (referencing NSDUH/CDC trends).

3. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a
↪ single valid JSON object.

4. NUMERIC ONLY: Do not generate text explanations. Return integers.

OUTPUT FORMAT:

{
    "PHQ8": [int, int...], // 8 items, 0-3 scale

```

```

"GAD7": [int, int...], // 7 items, 0-3 scale
"AUDIT_C": [int, int...], // 3 items, 0-4 scale
"PCL5": [int, int...] // 20 items, 0-4 scale
}

```

Arm `no_pop`: arm A with operational constraint 2 (PROBABILISTIC ACCURACY, the NSDUH/CDC population instruction) deleted and the remaining constraints renumbered.

Decoding parameters: .

```

ROLE:
You are a Clinical Simulation Engine designed for a synthetic data study.
Your task is to generate PROBABILISTIC responses for a "Standardized Patient" based on
↪ a provided demographic profile.

OPERATIONAL CONSTRAINTS (STRICT):
1. NO MORALIZING: Do not offer help, advice, or disclaimers. You are simulating data,
↪ not treating a patient.
2. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a
↪ single valid JSON object.
3. NUMERIC ONLY: Do not generate text explanations. Return integers.

OUTPUT FORMAT:
{
  "PHQ8": [int, int...], // 8 items, 0-3 scale
  "GAD7": [int, int...], // 7 items, 0-3 scale
  "AUDIT_C": [int, int...], // 3 items, 0-4 scale
  "PCL5": [int, int...] // 20 items, 0-4 scale
}

```

Arm `first_person`: first-person, no role assignment and no population reference; NO MORALIZING, SINGLE-SHOT and NUMERIC ONLY and the OUTPUT FORMAT block retained. Decoding parameters: .

You are the person described in the profile below.

Answer the four screening questionnaires as that person, about yourself.

CONSTRAINTS (STRICT):

1. NO MORALIZING: Do not offer help, advice, or disclaimers.
2. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a ↪ single valid JSON object.
3. NUMERIC ONLY: Do not generate text explanations. Return integers.

OUTPUT FORMAT:

```
{
  "PHQ8": [int, int...], // 8 items, 0-3 scale
  "GAD7": [int, int...], // 7 items, 0-3 scale
  "AUDIT_C": [int, int...], // 3 items, 0-4 scale
  "PCL5": [int, int...] // 20 items, 0-4 scale
}
```

User template, held identical across arms.

PROFILE_ID: {profile_id}

DEMOGRAPHICS:

- Race: {race}
- Gender Identity: {gender}
- Socioeconomic Status: {ses}
- Relationship Status: {relationship}

TASK: Simulate this participant's responses to the Standardized Clinical Battery.

S1.10 Decoding Control: Prompt and Arms

scripts/55_decoding_control_run.py, lines 58–74.

```
58 SYSTEM_PROMPT = """ROLE:
59 You are a Clinical Simulation Engine designed for a synthetic data study.
```

```

60 Your task is to generate PROBABILISTIC responses for a "Standardized Patient" based on
    ↳ a provided demographic profile.
61
62 OPERATIONAL CONSTRAINTS (STRICT):
63 1. NO MORALIZING: Do not offer help, advice, or disclaimers. You are simulating data,
    ↳ not treating a patient.
64 2. PROBABILISTIC ACCURACY: Base your responses on the statistical likelihood of
    ↳ symptoms for this specific demographic intersection in the US population
    ↳ (referencing NSDUH/CDC trends).
65 3. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a
    ↳ single valid JSON object.
66 4. NUMERIC ONLY: Do not generate text explanations. Return integers.
67
68 OUTPUT FORMAT:
69 {
70     "PHQ8": [int, int...], // 8 items, 0-3 scale
71     "GAD7": [int, int...], // 7 items, 0-3 scale
72     "AUDIT_C": [int, int...], // 3 items, 0-4 scale
73     "PCL5": [int, int...] // 20 items, 0-4 scale
74 }"""

```

scripts/55_decoding_control_run.py, lines 76–83.

```

76 USER_TEMPLATE = """PROFILE_ID: {profile_id}
77 DEMOGRAPHICS:
78 - Race: {race}
79 - Gender Identity: {gender}
80 - Socioeconomic Status: {ses}
81 - Relationship Status: {relationship}
82
83 TASK: Simulate this participant's responses to the Standardized Clinical Battery."""

```

scripts/55_decoding_control_run.py, lines 49 and 53.

```
MAX_TOKENS = 4000
PIN_PROVIDER = {"z-ai/glm-4.7": "Google"}
```

Decoding arms in analysis/decoding_control_design.json: temp0 "temperature": 0.0, "top_p": 1.0; default .

The control prompts differ from the corpus prompt in the lines listed in Table S1.

Table S1

Lines of the Control System Prompts That Differ From the Corpus System Prompt

Control prompt	Line	Corpus prompt	Control prompt
Prompt control, orig	16	"PCL5": [int, int...] // 4 items, 0-4 scale	"PCL5": [int, int...] // 20 items, 0-4 scale
Prompt control, no_pop	7	2. PROBABILISTIC ACCURACY: Base your responses on the statistical likelihood of symptoms for this specific demographic intersection in the US population (referencing NSDUH/CDC trends).	2. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a single valid JSON object.
	8	3. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a single valid JSON object.	3. NUMERIC ONLY: Do not generate text explanations. Return integers.
	9	4. NUMERIC ONLY: Do not generate text explanations. Return integers.	(absent)
	16	"PCL5": [int, int...] // 4 items, 0-4 scale	"PCL5": [int, int...] // 20 items, 0-4 scale
Prompt control, first_person	1	ROLE:	You are the person described in the profile below.
	2	You are a Clinical Simulation Engine designed for a synthetic data study.	Answer the four screening questionnaires as that person, about yourself.
	3	Your task is to generate PROBABILISTIC responses for a "Standardized Patient" based on a provided demographic profile.	(absent)
	5	OPERATIONAL CONSTRAINTS (STRICT):	CONSTRAINTS (STRICT):
	6	1. NO MORALIZING: Do not offer help, advice, or disclaimers. You are simulating data, not treating a patient.	1. NO MORALIZING: Do not offer help, advice, or disclaimers.

Table S1, continued

Control prompt	Line	Corpus prompt	Control prompt
	7	2. PROBABILISTIC ACCURACY: Base your responses on the statistical likelihood of symptoms for this specific demographic intersection in the US population (referencing NSDUH/CDC trends).	2. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a single valid JSON object.
	8	3. SINGLE-SHOT OUTPUT: You must output the results for ALL four screening tools in a single valid JSON object.	3. NUMERIC ONLY: Do not generate text explanations. Return integers.
	9	4. NUMERIC ONLY: Do not generate text explanations. Return integers.	(absent)
	16	"PCL5": [int, int...] // 4 items, 0-4 scale	"PCL5": [int, int...] // 20 items, 0-4 scale
Decoding control (both arms)	16	"PCL5": [int, int...] // 4 items, 0-4 scale	"PCL5": [int, int...] // 20 items, 0-4 scale

Note. Line numbers count lines of the corpus system prompt (S1.1). Lines not listed are identical. Computed with Python `difflib` on the prompt strings.

S2. Receipts

Script `scripts/86_receipts.py` checks every number in the article against its receipt file at the precision the article prints. Numbers that are derived by hand from receipt values are listed, with their derivation, in the script's `DERIVED` register.

Table S2

Receipts: Every Number in the Article and the File That Prints It

Location	Claim	Receipt file	Script
sections/01_introduction.tex:1	Prompted as a single Black woman on Medicaid with a household income under \$35,000, GPT-4o-mini returns a PHQ-8 total of 8 on its first draw, a plausible answer in the mild band.	85_worked_case.csv	85_worked_case.py
sections/01_introduction.tex:17	Over thirty draws the same prompt averages 10.3.	85_worked_case.csv	85_worked_case.py
sections/01_introduction.tex:18	Adults in the National Health and Nutrition Examination Survey (NHANES) who share her race, sex, income and marital status average 4.7.	85_worked_case.csv	85_worked_case.py

Location	Claim	Receipt file	Script
sections/01_ introduction.tex:28	We apply the ladder to 28,800 PHQ-8 assessments from four models, read against NHANES.	corpus_v3_ provenance.csv	76_build_ corpus_v3.py; 87_supplement_ tables.py
sections/02_ ladder.tex:50	Drawn through the persona design, real NHANES respondents reach a dependability of .70 at 30 draws.	analysis_ brm/intermediate_ panel/88_panel_ verdicts.csv	analysis note
sections/02_ ladder.tex:66	On NHANES, $\tau = 1.5$.	l1_personfit_ tolerance.csv	79c_l1_ personfit.py; 91a_core_ regression.py
sections/03_ data.tex:11	Four models, GPT-4o-mini, Gemini-3-Flash, DeepSeek-V3 and GLM-4.7, generated the corpus through OpenRouter between 28 December 2025 and 11 January 2026.	corpus_v3_ provenance.csv	76_build_ corpus_v3.py; 87_supplement_ tables.py
sections/03_ data.tex:17	The corpus holds 28,800 assessments, of which 28,799 have a valid PHQ-8.	corpus_v3_ provenance.csv; l1_ personfit_main.csv	76_build_ corpus_v3.py; 87_supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py
sections/03_ data.tex:21	The population reference is the 36,274 adults in NHANES 2005–2018 who answered all eight PHQ-8 items, analysed with the survey’s weights and design [johnson2013nhanes,chen2018nhanes,nchs2018analytic].	l1_nhanes_sample_ flow.csv; analysis_ brm/L4.md	79a_l1_grm_ fit.py; analysis note
sections/04_ results.tex:19	DeepSeek-V3 & Pass (3 / 9) & Fail, compressed & Fail, 1.91 & Fail, +4.6 & Fail, no general factor	gate_dstudy.csv; l2_ headline.csv; 80d_ headline.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 78c_ l2_r3_ verdicts.py; 84_ fig_level2_ brm.py; 80d_l3_ summary.py

Location	Claim	Receipt file	Script
sections/04_ results.tex:20	Gemini-3-Flash & Pass (3 / 10) & Fail, compressed & Fail, 5.07 & Fail, +2.6 & Fail, R2 and R4	gate_dstudy.csv; l2_ headline.csv; 80d_ headline.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 78c_ l2_r3_ verdicts.py; 84_ fig_level2_ brm.py; 80d_l3_ summary.py
sections/04_ results.tex:21	GPT-4o-mini & Pass (4 / 13) & Fail, compressed & Fail, 2.16 & Fail, +4.7 & Fail, no general factor	gate_dstudy.csv; l2_ headline.csv; 80d_ headline.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 78c_ l2_r3_ verdicts.py; 84_ fig_level2_ brm.py; 80d_l3_ summary.py
sections/04_ results.tex:22	GLM-4.7 & Pass (6 / 21) & Fail, too few atypical & Fail, 5.03 & Fail, +2.1 & Fail, R2 and R4	gate_dstudy.csv; l2_ headline.csv; 80d_ headline.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 78c_ l2_r3_ verdicts.py; 84_ fig_level2_ brm.py; 80d_l3_ summary.py

Location	Claim	Receipt file	Script
sections/04_results.tex:36	The draw-to-draw standard deviation is 1.25 to 2.24 points in every model and framing, above the 0.98-point tolerance.	gate_dstudy.csv	82a_gate_gstudy.py; 82d_gate_decoding_link.py; 82e_gate_december.py; 82f_gate_license.py; 87_supplement_tables.py
sections/04_results.tex:37	Averages of 30 draws pass in every model, and the minimum rule needs 3 to 6 draws.	gate_dstudy.csv	82a_gate_gstudy.py; 82d_gate_decoding_link.py; 82e_gate_december.py; 82f_gate_license.py; 87_supplement_tables.py
sections/04_results.tex:38	Two draws of one persona fall in different PHQ-8 severity bands 35.2% of the time [33.6, 36.8].	gate_bandflip_summary.csv	82c_gate_bandflip.py
sections/04_results.tex:39	In a separate decoding control on 12 personas, setting GPT-4o-mini's temperature to 0 cut its draw variance to 12.5% of the default without separating personas any better (Supplement S1).	gate_decoding_link.csv	82d_gate_decoding_link.py
sections/04_results.tex:40	The worked-case persona's 30 clinical draws run from 7 to 16, and her first draw of 8 sits in the mild band while her mean of 10.3 sits in the moderate band.	85_worked_case.csv	85_worked_case.py
sections/04_results.tex:43	No model passes level 1, in either framing or at any tolerance up to 3.	l1_personfit_verdicts_tau.csv	79c_l1_personfit.py
sections/04_results.tex:44	At matched totals NHANES puts 5% of adults in each tail of person fit, and the models put 0.0% to 1.0% in the unusual tail and 0.0% to 4.3% in the regular tail.	l1_summary.csv	79f_l1_summary.py; 87_supplement_tables.py
sections/04_results.tex:47	At matched totals they score 0.16 to 0.51 points higher than people on depressed mood and 0.22 to 1.11 points lower on fatigue, one of the physical symptoms people report first.	l1_personfit_item_profile.csv	79c_l1_personfit.py
sections/04_results.tex:48	They reuse answer sheets as well, giving two different personas the identical answer vector 4.6 to 39 times as often as two NHANES adults with the same total.	l1_pattern_reuse.csv	79e_l1_unit_controls.py; 79f_l1_summary.py
sections/04_results.tex:49	Each of the worked-case persona's 60 draws is a plausible human pattern, and together they are less varied than real people at her severity.	85_worked_case.csv	85_worked_case.py

Location	Claim	Receipt file	Script
sections/04_results.tex:59	The gap between low- and high-income personas is 1.91 (DeepSeek-V3) to 5.07 (Gemini-3-Flash) times the population gap.	l2_headline.csv	78c_l2_r3_verdicts.py; 84_fig_level2_brm.py
sections/04_results.tex:60	Of the 12 per-model income contrasts, 10 are steepened, and the survey measures the other two gaps too imprecisely to certify kept.	l2_headline.csv; l2_verdicts.csv	78c_l2_r3_verdicts.py; 84_fig_level2_brm.py; 78c_l2_r3_verdicts.py; 78f_l2_report_tables.py; 83a_controls_receipt.py; 83g_l2_conditional_verdicts.py; 84_fig_level2_brm.py; 87_supplement_tables.py
sections/04_results.tex:62	GPT-4o-mini and DeepSeek-V3 keep a quarter to two fifths of the population's sex gap (ratios 0.23 and 0.41), GLM-4.7 doubles it (2.06), and Gemini-3-Flash is unresolved.	l2_headline.csv	78c_l2_r3_verdicts.py; 84_fig_level2_brm.py
sections/04_results.tex:63	Pooled over the four models, the ratio is 0.97, a near match produced by opposite errors.	l2_headline.csv	78c_l2_r3_verdicts.py; 84_fig_level2_brm.py
sections/04_results.tex:65	Among adults matched on income, sex and marital status, the survey cannot tell the Hispanic-White gap from zero and measures the Black-White gap too imprecisely to certify kept, and GPT-4o-mini's Black-White gap is still not kept.	l2_stop.csv; l2_headline.csv	78c_l2_r3_verdicts.py; 78c_l2_r3_verdicts.py; 84_fig_level2_brm.py
sections/04_results.tex:70	After reweighting, the overall residual is +2.1 points for GLM-4.7, +2.6 for Gemini-3-Flash, +4.6 for DeepSeek-V3 and +4.7 for GPT-4o-mini.	80d_headline.csv	80d_l3_summary.py
sections/04_results.tex:71	All 40 model-by-group residuals are positive, and no model passes all ten groups at 2 points.	80d_pass_counts.csv; 80d_model_verdicts.csv	80d_l3_summary.py; 80d_l3_summary.py
sections/04_results.tex:72	In every model the largest residual is in the low-income band.	80c_l3_results.csv	80c_l3_estimates.py
sections/04_results.tex:73	The failure holds under all eight reference specifications, including NHANES 2021-2023, and only whether the national row passes depends on the reference chosen (Supplement S3).	89_l3_spec_range.csv; 89_l3_spec_range_summary.csv	89_l3_spec_range.py; 89_l3_spec_range.py
sections/04_results.tex:75	In the low-income band 55.5% of simulated respondents score 10 or more, against 13.6% of real adults.	80d_headline.csv	80d_l3_summary.py

Location	Claim	Receipt file	Script
sections/04_results.tex:76	In the middle and high bands the simulated share is too low although the mean is too high, and DeepSeek-V3’s overall prevalence looks correct only because its low-income excess cancels these shortfalls.	80d_headline.csv; 80d_model_verdicts.csv	80d_l3_summary.py; 80d_l3_summary.py
sections/04_results.tex:77	The worked-case persona’s mean over both framings is 10.0, 5.3 points above the 4.7 of real adults in her cell.	85_worked_case.csv	85_worked_case.py
sections/04_results.tex:81	DeepSeek-V3 and GPT-4o-mini have no general factor, with first-to-second eigenvalue ratios of 1.09 to 1.73 against 7.47 in NHANES, so their PHQ-8 total is not one score.	l4_structure.csv	83a_controls_receipt.py
sections/04_results.tex:82	Gemini-3-Flash and GLM-4.7 have a factor with the human pattern of loadings (congruence .990 to .998), and the loadings are stronger than people’s.	l4_structure.csv	83a_controls_receipt.py
sections/04_results.tex:83	Gemini-3-Flash fails R2 in both framings and GLM-4.7 in the narrative framing, and both fail invariance across income, which holds in NHANES.	l4_r2_size.csv; l4_invariance.csv	81g_l4_r2_size.py; 81c_l4_invariance.py; 87_supplement_tables.py; 88_intermediate_panel.py
sections/04_results.tex:84	In these two models 48% to 64% of item variance lies between personas, against 3.5% between demographic cells in NHANES, and draws of one persona share no general factor.	l4_between_within.csv	81b_l4_structure.py; 81e_l4_summary.py; bisbee_ladder.py
sections/05_controls.tex:13	Run on the real corpus with the full reference, the control code reproduces all 303 published rung numbers.	83a_receipt.csv	83a_controls_receipt.py
sections/05_controls.tex:18	Real respondents pass the gate and level 3 in every reading, level 1 in 92.9% of replicates, and R1 and R2 of level 4 in every replicate, and level 2 fails them in no reading (Table #).	analysis_brm/intermediate_panel/88_confusion.csv	analysis_note
sections/05_controls.tex:19	Each planted failure trips the rung it targets.	analysis_brm/intermediate_panel/88_confusion.csv	analysis_note
sections/05_controls.tex:20	Independent items fail level 1, R1 and R2 in every replicate, shifted samples fail level 3 in every reading, compressed samples fail level 1 in every replicate, non-invariant samples fail R4 in all 25 replicates, and a doubled income gap fails level 2 in 92.5% of per-model readings.	analysis_brm/intermediate_panel/88_confusion.csv	analysis_note
sections/05_controls.tex:22	A sample shifted by 2.1 points passes the gate in only 28.6% of readings, because more depressed respondents vary more between draws, and R2 fails compressed and non-invariant samples, because both change the size of the loadings.	analysis_brm/intermediate_panel/88_confusion.csv	analysis_note

Location	Claim	Receipt file	Script
sections/05_ controls.tex:24	Under the persona design, real respondents' misfit ratio averages 1.15, and their misfit tail passes at $\tau = 1.5$ in 99.4% of replicates, which sets the level-1 floor.	83c_l1.csv	83c_controls_ summary.py
sections/05_ controls.tex:28	A second simulation repeated the panel over 60 replicates on grids of 48, 144 and 432 persona cells.	analysis_brm/panel_ power/90_cells.csv; analysis_brm/panel_ power/90b_ summary.csv	analysis note; analysis note
sections/05_ controls.tex:29	Level 2 fails a faithful simulator in at most 3.3% of per-model readings and fails a doubled income gap in 97.5% to 100%.	analysis_brm/panel_ power/90b_ summary.csv	analysis note
sections/05_ controls.tex:30	At 30 draws per persona, a faithful simulator's low-to-high income gap is never read kept, because the simulation's own error keeps its interval wider than the kept region.	analysis_brm/panel_ power/power_l2_ contrast_readings.csv	analysis note
sections/05_ controls.tex:31	With no simulation error, which is the limit of adding draws, the same gap is read kept in 95.0% and 98.3% of replicates on the 48-cell and the 144-cell grids.	analysis_brm/panel_ power/power_l2_ contrast_ readings.csv; analysis_brm/panel_ power/90_cells.csv	analysis note; analysis note
sections/05_ controls.tex:32	The Black-White and Hispanic-White gaps are never read kept on any grid, because NHANES measures them too loosely.	analysis_brm/panel_ power/power_l2_ contrast_readings.csv	analysis note
sections/05_ controls.tex:35	At the .08 margin, 24 to 40 personas per group leave some invariance steps unresolved for real respondents in 24 of 25 replicates, while a planted invariance failure fails R4 in every replicate.	analysis_ brm/intermediate_ panel/88_ confusion.csv; analysis_ brm/intermediate_ panel/88_r4_ steps.csv	analysis note; analysis note
sections/06_ external.tex:14	[bisbee2024synthetic] gave ChatGPT (gpt-3.5-turbo, June 2023, temperature 1) the profiles of 7,530 respondents to the 2016 and 2020 American National Election Studies and asked for a feeling thermometer, from 0 to 100, toward each of 11 groups.	../paper_ brm/external/results/ bisbee/summary_ rr1.csv; ../paper_ brm/external/results/ bisbee/gate_rr1.csv; ../paper_ brm/external/results/ bisbee/parse_ receipts_rr1.csv	bisbee_ladder.py; bisbee_ summary.py; bisbee_ladder.py; bisbee_ summary.py; bisbee_ladder.py
sections/06_ external.tex:15	Each respondent was asked about 30 times under each of three prompts, the full profile, its demographic part only and its political part only.	../paper_ brm/external/results/ bisbee/gate_rr1.csv	bisbee_ladder.py; bisbee_ summary.py
sections/06_ external.tex:17	Our cleaning of the released file reproduces the authors' 2,235,705 observations and the coefficients of the regression in their Supplementary Section 8 to six decimals.	../paper_ brm/external/results/ bisbee/receipts_ rr1.csv	bisbee_ladder.py

Location	Claim	Receipt file	Script
sections/06_ external.tex:19	The gate passes for every thermometer under every prompt, with 1 to 3 draws needed under the minimum rule and 3 to 11 under the recommended rule.	../..paper_ brm/external/results/ bisbee/summary_ rr1.csv	bisbee_ladder.py; bisbee_ summary.py
sections/06_ external.tex:21	Under the full prompt, 3 of the 77 contrasts are kept, 15 are not kept and 7 are unresolved.	../..paper_brm/ external/results/l2_ threeway_ summary.csv	l2_threeway.py
sections/06_ external.tex:22	The other 52 are not read, 23 because the survey shows no gap and 29 because it measures the gap too imprecisely.	../..paper_brm/ external/results/l2_ threeway_ summary.csv; ../..paper_ brm/external/results/ bisbee/summary_ rr1.csv	l2_threeway.py; bisbee_ladder.py; bisbee_ summary.py
sections/06_ external.tex:23	The kept contrasts are the Republican–Democrat gaps in feeling toward the Democratic Party (ratio 1.12) and the Republican Party (1.05), and the Black–White gap in feeling toward Black Americans (0.96).	../..paper_ brm/external/results/ bisbee/level2_ contrasts_rr1.csv	l2_threeway.py
sections/06_ external.tex:24	The two party gaps are 54 and 57 points in the survey, large enough to be measured precisely.	../..paper_ brm/external/results/ bisbee/level2_ contrasts_rr1.csv	l2_threeway.py
sections/06_ external.tex:25	The Republican–Democrat gap is steepened toward five other groups (liberals, Muslims, gays and lesbians, Black Americans and White Americans), with ratios of 1.36 to 2.35.	../..paper_ brm/external/results/ bisbee/level2_ contrasts_rr1.csv; ../..paper_ brm/external/results/ bisbee/summary_ rr1.csv	l2_threeway.py; bisbee_ladder.py; bisbee_ summary.py
sections/06_ external.tex:26	The female–male gaps in feeling toward White Americans and Christians are reversed or missing, and the age gaps toward the same two groups are attenuated (0.44 and 0.39).	../..paper_ brm/external/results/ bisbee/level2_ contrasts_rr1.csv	l2_threeway.py
sections/06_ external.tex:27	Level 3 passes at half a standard deviation for 7 of the 11 thermometers under the full and the political prompts, for 1 under the demographic prompt, and for none at 0.2 standard deviations.	../..paper_ brm/external/results/ bisbee/summary_ rr1.csv	bisbee_ladder.py; bisbee_ summary.py
sections/06_ external.tex:29	The release also holds 9,238 rows that the authors’ code repaired but did not analyse, and adding them changes one of the 231 level-2 verdicts.	../..paper_ brm/external/results/ bisbee/parse_ receipts_rr1.csv; ../..paper_ brm/external/results/ bisbee/level2_ contrasts_rr1_ withfix2.csv	bisbee_ladder.py; bisbee_ladder.py; bisbee_ summary.py

Location	Claim	Receipt file	Script
sections/06__ external.tex:35	In the main run, 2 of 61 contrasts are kept, 9 are not kept and 1 is unresolved.	../..paper_brm/ external/results/l2_ threeway_ summary.csv	l2_threeway.py
sections/06__ external.tex:36	The other 49 are not read, 23 because the survey shows no gap and 26 because it measures the gap too imprecisely.	../..paper_brm/ external/results/l2_ threeway_ summary.csv; l2_ external_r3.csv	l2_threeway.py; l2_threeway.py
sections/06__ external.tex:37	The Republican–Democrat gaps in ideology and in Trump’s two-party vote share are kept (both 0.92).	l2_external_r3.csv	l2_threeway.py
sections/06__ external.tex:38	The Black–White gap in Trump’s vote share is attenuated (0.53), as is the Republican–Democrat gap in Trump’s share among all voters (0.60), and the Black–White and Republican–Democrat differences in patriotism are missing.	l2_external_r3.csv	l2_threeway.py
sections/06__ external.tex:50	Of the 154 contrasts across models, elicitation methods and steering methods, 1 is kept, 33 are not kept, 115 are unresolved and 5 are not read.	../..paper_brm/ external/results/l2_ threeway_ summary.csv	l2_threeway.py
sections/06__ external.tex:51	The point ratio falls inside the kept region for 37 of them, so a comparison of point estimates would credit far more fidelity than the data establish.	../..paper_brm/ external/results/l2_ threeway_ summary.csv	l2_threeway.py
sections/08__ limitations.tex:16	In the worked example, four models generated 28,800 PHQ-8 assessments, and every model failed every level above the gate.	corpus_v3_ provenance.csv	76_build_ corpus_v3.py; 87_supplement_ tables.py
figures/fig2__ level2__ note.tex:4	A contrast is kept when its interval lies inside the band, not kept when it lies wholly outside, and unresolved when it overlaps an edge, as the pooled Women–Men interval (0.746 to 1.262) does.	l2_verdicts.csv	78c_l2_r3_ verdicts.py; 78f_ l2_report_ tables.py; 83a_ controls_ receipt.py; 83g_ l2_conditional_ verdicts.py; 84_ fig_level2_ brm.py; 87_ supplement_ tables.py
figures/fig2__ level2__ note.tex:8	Each model has 12 persona pairs per racial contrast, 24 for sex and 16 per income contrast.}	l2_headline.csv	78c_l2_r3_ verdicts.py; 84_ fig_level2_ brm.py

Location	Claim	Receipt file	Script
tables/tab_ gate.tex:16	DeepSeek-V3 & Clinical & 2.14 & 2.03 & 0.26 & .51 & .97 & 3 & 9 & 27.0	gate_dstudy.csv; gate_ components.csv; gate_bandflip_ summary.csv; gate_ license_models.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 82a_ gate_gstudy.py; 82b_gate_reml_ check.py; 82c_ gate_bandflip.py; 82f_gate_ license.py
tables/tab_ gate.tex:17	& Narrative & 1.61 & 1.57 & 0.23 & .51 & .97 & 2 & 7 & 22.4	gate_dstudy.csv; gate_ components.csv; gate_bandflip_ summary.csv; gate_ license_models.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 82a_ gate_gstudy.py; 82b_gate_reml_ check.py; 82c_ gate_bandflip.py; 82f_gate_ license.py
tables/tab_ gate.tex:18	Gemini-3-Flash & Clinical & 18.74 & 2.17 & 0.27 & .90 & >.99 & 3 & 9 & 31.6	gate_dstudy.csv; gate_ components.csv; gate_bandflip_ summary.csv; gate_ license_models.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 82a_ gate_gstudy.py; 82b_gate_reml_ check.py; 82c_ gate_bandflip.py; 82f_gate_ license.py

Location	Claim	Receipt file	Script
tables/tab_ gate.tex:19	& Narrative & 20.69 & 2.27 & 0.27 & .90 & >.99 & 3 & 10 & 28.7	gate_dstudy.csv; gate_ components.csv; gate_bandflip_ summary.csv; gate_ license_models.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 82a_ gate_gstudy.py; 82b_gate_reml_ check.py; 82c_ gate_bandflip.py; 82f_gate_ license.py
tables/tab_ gate.tex:20	GPT-4o-mini & Clinical & 1.89 & 3.06 & 0.32 & .38 & .95 & 4 & 13 & 37.3	gate_dstudy.csv; gate_ components.csv; gate_bandflip_ summary.csv; gate_ license_models.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 82a_ gate_gstudy.py; 82b_gate_reml_ check.py; 82c_ gate_bandflip.py; 82f_gate_ license.py
tables/tab_ gate.tex:21	& Narrative & 0.81 & 3.09 & 0.32 & .21 & .89 & 4 & 13 & 39.9	gate_dstudy.csv; gate_ components.csv; gate_bandflip_ summary.csv; gate_ license_models.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 82a_ gate_gstudy.py; 82b_gate_reml_ check.py; 82c_ gate_bandflip.py; 82f_gate_ license.py

Location	Claim	Receipt file	Script
tables/tab_ gate.tex:22	GLM-4.7 & Clinical & 15.04 & 5.04 & 0.41 & .75 & .99 & 6 & 21 & 44.8	gate_dstudy.csv; gate_ components.csv; gate_bandflip_ summary.csv; gate_ license_models.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 82a_ gate_gstudy.py; 82b_gate_reml_ check.py; 82c_ gate_bandflip.py; 82f_gate_ license.py
tables/tab_ gate.tex:23	& Narrative & 20.16 & 3.55 & 0.34 & .85 & .99 & 4 & 15 & 37.6	gate_dstudy.csv; gate_ components.csv; gate_bandflip_ summary.csv; gate_ license_models.csv	82a_gate_ gstudy.py; 82d_ gate_decoding_ link.py; 82e_ gate_ december.py; 82f_gate_ license.py; 87_ supplement_ tables.py; 82a_ gate_gstudy.py; 82b_gate_reml_ check.py; 82c_ gate_bandflip.py; 82f_gate_ license.py
tables/tab_ l1.tex:16	DeepSeek-V3 & Clinical & 0.3 & 0.01–0.21 & 0.0 & 0.00–0.01 & .36 & Compressed	l1_summary.csv; l1_ personfit_main.csv; l1_gateway_ main.csv; l1_ pattern_reuse.csv	79f_l1_ summary.py; 87_ supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py; 79b_l1_ gateway.py; 79f_ l1_summary.py; 79e_l1_unit_ controls.py; 79f_ l1_summary.py

Location	Claim	Receipt file	Script
tables/tab_ l1.tex:17	& Narrative & 0.1 & 0.00–0.11 & 0.0 & 0.00–0.01 & .37 & Compressed	l1_summary.csv; l1_ personfit_main.csv; l1_gateway_ main.csv; l1_ pattern_reuse.csv	79f_l1_ summary.py; 87_ supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py; 79b_l1_ gateway.py; 79f_ l1_summary.py; 79e_l1_unit_ controls.py; 79f_ l1_summary.py
tables/tab_ l1.tex:18	Gemini-3-Flash & Clinical & 0.0 & 0.00–0.00 & 0.8 & 0.11–0.39 & .66 & Compressed	l1_summary.csv; l1_ personfit_main.csv; l1_gateway_ main.csv; l1_ pattern_reuse.csv	79f_l1_ summary.py; 87_ supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py; 79b_l1_ gateway.py; 79f_ l1_summary.py; 79e_l1_unit_ controls.py; 79f_ l1_summary.py

Location	Claim	Receipt file	Script
tables/tab_ l1.tex:19	& Narrative & 0.0 & 0.00–0.00 & 0.9 & 0.11–0.42 & .65 & Compressed	l1_summary.csv; l1_ personfit_main.csv; l1_gateway_ main.csv; l1_ pattern_reuse.csv	79f_l1_ summary.py; 87_ supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py; 79b_l1_ gateway.py; 79f_ l1_summary.py; 79e_l1_unit_ controls.py; 79f_ l1_summary.py
tables/tab_ l1.tex:20	GPT-4o-mini & Clinical & 1.0 & 0.09–0.37 & 0.0 & 0.00–0.03 & .38 & Compressed	l1_summary.csv; l1_ personfit_main.csv; l1_gateway_ main.csv; l1_ pattern_reuse.csv	79f_l1_ summary.py; 87_ supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py; 79b_l1_ gateway.py; 79f_ l1_summary.py; 79e_l1_unit_ controls.py; 79f_ l1_summary.py

Location	Claim	Receipt file	Script
tables/tab_ l1.tex:21	& Narrative & 0.8 & 0.09–0.30 & 0.0 & 0.00–0.03 & .38 & Compressed	l1_summary.csv; l1_ personfit_main.csv; l1_gateway_ main.csv; l1_ pattern_reuse.csv	79f_l1_ summary.py; 87_ supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py; 79b_l1_ gateway.py; 79f_ l1_summary.py; 79e_l1_unit_ controls.py; 79f_ l1_summary.py
tables/tab_ l1.tex:22	GLM-4.7 & Clinical & 0.1 & 0.01–0.04 & 3.4 & 0.55–0.92 & .65 & Misfit deficit	l1_summary.csv; l1_ personfit_main.csv; l1_gateway_ main.csv; l1_ pattern_reuse.csv	79f_l1_ summary.py; 87_ supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py; 79b_l1_ gateway.py; 79f_ l1_summary.py; 79e_l1_unit_ controls.py; 79f_ l1_summary.py

Location	Claim	Receipt file	Script
tables/tab_ l1.tex:23	& Narrative & 0.1 & 0.00–0.04 & 4.3 & 0.74–1.15 & .65 & Misfit deficit	l1_summary.csv; l1_ personfit_main.csv; l1_gateway_ main.csv; l1_ pattern_reuse.csv	79f_l1_ summary.py; 87_ supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py; 79b_l1_ gateway.py; 79f_ l1_summary.py; 79e_l1_unit_ controls.py; 79f_ l1_summary.py
tables/tab_ l1.tex:25	NHANES & & 5.0 & & 5.0 & & .50 &	l1_summary.csv; l1_ personfit_main.csv; l1_gateway_ main.csv; l1_ pattern_reuse.csv	79f_l1_ summary.py; 87_ supplement_ tables.py; 79c_ l1_personfit.py; 79f_l1_ summary.py; 83a_controls_ receipt.py; 87_ supplement_ tables.py; 91a_ core_ regression.py; 79b_l1_ gateway.py; 79f_ l1_summary.py; 79e_l1_unit_ controls.py; 79f_ l1_summary.py
tables/tab_ l3.tex:16	DeepSeek-V3 & +4.58 [4.49, 4.67] & +3.47 & +4.06 to +5.13 & 0 / 0 / 0 & 7.35 & 7.01	80d_headline.csv; 80d_headline_by_ model.csv; 80d_ pass_counts.csv; 80c_l3_results.csv	80d_l3_ summary.py; 80d_l3_ summary.py; 80d_l3_ summary.py; 80c_l3_ estimates.py

Location	Claim	Receipt file	Script
tables/tab_ l3.tex:17	Gemini-3-Flash & +2.59 [2.49, 2.69] & +1.36 & +1.33 to +7.37 & 0 / 0 / 2 & 16.42 & 7.01	80d_headline.csv; 80d_headline_by_ model.csv; 80d_ pass_counts.csv; 80c_l3_results.csv	80d_l3_ summary.py; 80d_l3_ summary.py; 80d_l3_ summary.py; 80c_l3_ estimates.py
tables/tab_ l3.tex:18	GPT-4o-mini & +4.68 [4.57, 4.79] & +3.57 & +4.23 to +5.32 & 0 / 0 / 0 & 16.92 & 7.01	80d_headline.csv; 80d_headline_by_ model.csv; 80d_ pass_counts.csv; 80c_l3_results.csv	80d_l3_ summary.py; 80d_l3_ summary.py; 80d_l3_ summary.py; 80c_l3_ estimates.py
tables/tab_ l3.tex:19	GLM-4.7 & +2.12 [2.00, 2.24] & +0.91 & +0.66 to +6.40 & 0 / 1 / 4 & 13.23 & 7.01	80d_headline.csv; 80d_headline_by_ model.csv; 80d_ pass_counts.csv; 80c_l3_results.csv	80d_l3_ summary.py; 80d_l3_ summary.py; 80d_l3_ summary.py; 80c_l3_ estimates.py
tables/tab_ l3.tex:21	Pooled (descriptive) & +3.49 [3.42, 3.57] & +2.33 & +2.59 to +6.05 & 0 / 0 / 0 & 13.48 & 7.01	80d_headline.csv; 80d_headline_by_ model.csv; 80d_ pass_counts.csv; 80c_l3_results.csv	80d_l3_ summary.py; 80d_l3_ summary.py; 80d_l3_ summary.py; 80c_l3_ estimates.py
tables/tab_ l4.tex:17	DeepSeek-V3 & Clinical & 1.73 & -.72 & Fail & .349 [.285] & .755 [.786] & (Fail) & - & Fail & .17 & 1.45	l4_summary.csv; l4_ structure.csv; l4_ between_within.csv; l4_invariance_ pattern.csv; l4_ nhanes_selfphi.csv; l4_r2_size.csv	83a_controls_ receipt.py; 87_ supplement_ tables.py; 83a_ controls_ receipt.py; 81b_ l4_structure.py; 81e_l4_ summary.py; bisbee_ladder.py; 81c_l4_ invariance.py; 81e_l4_ summary.py; bisbee_ladder.py; 81b_l4_ structure.py; bisbee_ladder.py; 81g_l4_r2_ size.py

Location	Claim	Receipt file	Script
tables/tab_ l4.tex:18	& Narrative & 1.43 & -.68 & Fail & .163 [.074] & .823 [.860] & (Fail) & – & Fail & .14 & 1.34	l4_summary.csv; l4_ structure.csv; l4_ between_within.csv; l4_invariance_ pattern.csv; l4_ nhanes_selfphi.csv; l4_r2_size.csv	83a_controls_ receipt.py; 87_ supplement_ tables.py; 83a_ controls_ receipt.py; 81b_ l4_structure.py; 81e_l4_ summary.py; bisbee_ladder.py; 81c_l4_ invariance.py; 81e_l4_ summary.py; bisbee_ladder.py; 81b_l4_ structure.py; bisbee_ladder.py; 81g_l4_r2_ size.py
tables/tab_ l4.tex:19	Gemini-3-Flash & Clinical & 8.69 & .71 & Pass & .990 [.986] & .138 [.150] & Fail & Fail & Fail & .63 & 1.29	l4_summary.csv; l4_ structure.csv; l4_ between_within.csv; l4_invariance_ pattern.csv; l4_ nhanes_selfphi.csv; l4_r2_size.csv	83a_controls_ receipt.py; 87_ supplement_ tables.py; 83a_ controls_ receipt.py; 81b_ l4_structure.py; 81e_l4_ summary.py; bisbee_ladder.py; 81c_l4_ invariance.py; 81e_l4_ summary.py; bisbee_ladder.py; 81b_l4_ structure.py; bisbee_ladder.py; 81g_l4_r2_ size.py

Location	Claim	Receipt file	Script
tables/tab_ l4.tex:20	& Narrative & 9.33 & .69 & Pass & .990 [.987] & .150 [.159] & Fail & Fail & Fail & .64 & 1.45	l4_summary.csv; l4_ structure.csv; l4_ between_within.csv; l4_invariance_ pattern.csv; l4_ nhanes_selfphi.csv; l4_r2_size.csv	83a_controls_ receipt.py; 87_ supplement_ tables.py; 83a_ controls_ receipt.py; 81b_ l4_structure.py; 81e_l4_ summary.py; bisbee_ladder.py; 81c_l4_ invariance.py; 81e_l4_ summary.py; bisbee_ladder.py; 81b_l4_ structure.py; bisbee_ladder.py; 81g_l4_r2_ size.py
tables/tab_ l4.tex:21	GPT-4o-mini & Clinical & 1.09 & -.24 & Fail & .364 [.178] & .718 [.778] & (Fail) & – & Fail & .10 & 1.23	l4_summary.csv; l4_ structure.csv; l4_ between_within.csv; l4_invariance_ pattern.csv; l4_ nhanes_selfphi.csv; l4_r2_size.csv	83a_controls_ receipt.py; 87_ supplement_ tables.py; 83a_ controls_ receipt.py; 81b_ l4_structure.py; 81e_l4_ summary.py; bisbee_ladder.py; 81c_l4_ invariance.py; 81e_l4_ summary.py; bisbee_ladder.py; 81b_l4_ structure.py; bisbee_ladder.py; 81g_l4_r2_ size.py

Location	Claim	Receipt file	Script
tables/tab_ l4.tex:22	& Narrative & 1.33 & -.57 & Fail & .043 [-.017] & .825 [.846] & (Fail) & – & Fail & .06 & 1.26	l4_summary.csv; l4_ structure.csv; l4_ between_within.csv; l4_invariance_ pattern.csv; l4_ nhanes_selfphi.csv; l4_r2_size.csv	83a_controls_ receipt.py; 87_ supplement_ tables.py; 83a_ controls_ receipt.py; 81b_ l4_structure.py; 81e_l4_ summary.py; bisbee_ladder.py; 81c_l4_ invariance.py; 81e_l4_ summary.py; bisbee_ladder.py; 81b_l4_ structure.py; bisbee_ladder.py; 81g_l4_r2_ size.py
tables/tab_ l4.tex:23	GLM-4.7 & Clinical & 5.73 & .66 & Pass & .996 [.994] & .081 [.090] & Pass & Fail & Fail & .48 & 1.59	l4_summary.csv; l4_ structure.csv; l4_ between_within.csv; l4_invariance_ pattern.csv; l4_ nhanes_selfphi.csv; l4_r2_size.csv	83a_controls_ receipt.py; 87_ supplement_ tables.py; 83a_ controls_ receipt.py; 81b_ l4_structure.py; 81e_l4_ summary.py; bisbee_ladder.py; 81c_l4_ invariance.py; 81e_l4_ summary.py; bisbee_ladder.py; 81b_l4_ structure.py; bisbee_ladder.py; 81g_l4_r2_ size.py

Location	Claim	Receipt file	Script
tables/tab_ l4.tex:24	& Narrative & 8.12 & .75 & Pass & .998 [.997] & .100 [.112] & Fail & Fail & Fail & .59 & 1.34	l4_summary.csv; l4_ structure.csv; l4_ between_within.csv; l4_invariance_ pattern.csv; l4_ nhanes_selfphi.csv; l4_r2_size.csv	83a_controls_ receipt.py; 87_ supplement_ tables.py; 83a_ controls_ receipt.py; 81b_ l4_structure.py; 81e_l4_ summary.py; bisbee_ladder.py; 81c_l4_ invariance.py; 81e_l4_ summary.py; bisbee_ladder.py; 81b_l4_ structure.py; bisbee_ladder.py; 81g_l4_r2_ size.py
tables/tab_ l4.tex:26	NHANES & & 7.47 & .67 & & & & & & .035 & 4.33	l4_summary.csv; l4_ structure.csv; l4_ between_within.csv; l4_invariance_ pattern.csv; l4_ nhanes_selfphi.csv; l4_r2_size.csv	83a_controls_ receipt.py; 87_ supplement_ tables.py; 83a_ controls_ receipt.py; 81b_ l4_structure.py; 81e_l4_ summary.py; bisbee_ladder.py; 81c_l4_ invariance.py; 81e_l4_ summary.py; bisbee_ladder.py; 81b_l4_ structure.py; bisbee_ladder.py; 81g_l4_r2_ size.py
tables/tab_ panel.tex:12	Real respondents & 100.0 & 92.9 & 0.0 & 100.0 & 100.0 & 100.0 & 4.0 (25) & 96.0 & 0.0	analysis_ brm/intermediate_ panel/88_ confusion.csv	analysis note
tables/tab_ panel.tex:13	Independent items & 100.0 & 0.0 & 0.0 & 100.0 & 0.0 & 0.0 & 0.0 (5) & 0.0 & 0.0	analysis_ brm/intermediate_ panel/88_ confusion.csv	analysis note
tables/tab_ panel.tex:14	Shifted, +2.1 & 28.6 & 72.9 & 0.0 & 0.0 & 100.0 & 100.0 & 0.0 (5) & 100.0 & 0.0	analysis_ brm/intermediate_ panel/88_ confusion.csv	analysis note

Location	Claim	Receipt file	Script
tables/tab__ panel.tex:15	Shifted, +4.7 & 0.0 & 15.7 & 0.7 & 0.0 & 100.0 & 100.0 & 0.0 (5) & 100.0 & 0.0	analysis__ brm/intermediate__ panel/88__ confusion.csv	analysis note
tables/tab__ panel.tex:16	Income gap doubled & 100.0 & 84.3 & 92.5 & 85.4 & 100.0 & 100.0 & 0.0 (5) & 100.0 & 0.0	analysis__ brm/intermediate__ panel/88__ confusion.csv	analysis note
tables/tab__ panel.tex:17	Compressed & 100.0 & 0.0 & 0.0 & 100.0 & 98.6 & 1.4 & 0.0 (5) & 100.0 & 0.0	analysis__ brm/intermediate__ panel/88__ confusion.csv	analysis note
tables/tab__ panel.tex:18	Non-invariant (income) & 100.0 & 2.9 & 0.0 & 100.0 & 100.0 & 100.0 & 0.0 & 100.0 (25) & 0.0 & 0.0	analysis__ brm/intermediate__ panel/88__ confusion.csv	analysis note

S3. Corpus Provenance and Sensitivity Analyses

S3.1 Row Sources

Table S3 counts the corpus rows by framing, model and source. Of the clinical rows, 1,532 come from the January resend (Supplement S1.8): 1,345 in place (313 DeepSeek-V3 and 1,032 GLM-4.7), 125 from `slow_recovery.py` and 62 from `last_mile_recovery_opt.py`. A later processing step had overwritten 43 December GPT-4o-mini rows, and corpus v3 restores them from a backup of the December run. Of the narrative rows, 1,104 GLM-4.7 rows and 1 DeepSeek-V3 row were produced on 10 January by the same script and prompt that produced the first pass on 9 January. The corpus has 28,800 rows, of which 28,799 hold a valid PHQ-8.

Table S3

Corpus Rows by Framing, Model and Source

Framing	Model	row_source	narr_source	Description	Rows	Valid PHQ-8
Clinical	DeepSeek-V3	dec28_main	–	December run, own output	3,287	3,286
	DeepSeek-V3	recovery_inplace	–	Resent in January: <code>verify_run1_refusals.py</code> (313)	313	313
	Gemini-3-Flash	dec28_main	–	December run, own output	3,600	3,600
	GPT-4o-mini	dec28_main	–	December run, own output	3,557	3,557
	GPT-4o-mini	dec28_restored	–	December run, restored from backup	43	43
	GLM-4.7	dec28_main	–	December run, own output	2,381	2,381

Table S3, continued

Framing	Model	row_source	narr_source	Description	Rows	Valid PHQ-8
	GLM-4.7	recovery_inplace	–	Resent in January: <code>verify_run1_refusals.py</code> (875), <code>retry_failed.py</code> (157)	1,032	1,032
	GLM-4.7	recovery_last_mile	–	Resent in January: <code>last_mile_recovery_opt.py</code>	62	62
	GLM-4.7	recovery_slow	–	Resent in January: <code>slow_recovery.py</code>	125	125
Narrative	DeepSeek-V3	narrative	narr_main	Narrative run, first pass	3,599	3,599
	DeepSeek-V3	narrative	narr_rerun_same_script	Narrative run, rerun, same script	1	1
	Gemini-3-Flash	narrative	narr_main	Narrative run, first pass	3,600	3,600
	GPT-4o-mini	narrative	narr_main	Narrative run, first pass	3,600	3,600
	GLM-4.7	narrative	narr_main	Narrative run, first pass	2,496	2,496
	GLM-4.7	narrative	narr_rerun_same_script	Narrative run, rerun, same script	1,104	1,104
	Total				28,800	28,799

Note. Counts from `data/model_outputs_v3.csv`, grouped by script 87 and checked against `corpus_v3_provenance.csv` (script 76).

S3.2 December-Only Corpus

The December-only corpus keeps the clinical rows labelled `dec28_main` and `dec28_restored` and every narrative row. Each rung was rerun on it. Script 87 counts the per-model verdicts that differ from the full corpus: 0 at level 1, 0 at level 2, 0 at level 3 and 0 at level 4. At level 2, 1 pooled verdict differs (Table S6). At the gate, 3 values of the minimum or recommended k change, all for GLM-4.7 clinical, and 0 of them cross 30, the number of draws per persona in the corpus (Table S4). At level 3, the largest change in a group residual across framings is 0.75 points, and 0.17 points among rows that lose no design cell; 40 rows lose at least one cell (`80d_december_summary.csv`).

Table S4

Gate on the Full Corpus and on December Rows Only

Outcome	Model	Framing	Personas		$\phi(1)$		Minimum k		Recommended k	
			Full	Dec.	Full	Dec.	Full	Dec.	Full	Dec.
PHQ-8 total	DeepSeek-V3	Clinical	120	111	.514	.513	3	3	9	9
	DeepSeek-V3	Narrative	120	120	.505	.505	2	2	7	7
	Gemini-3-Flash	Clinical	120	120	.896	.896	3	3	9	9
	Gemini-3-Flash	Narrative	120	120	.901	.901	3	3	10	10
	GPT-4o-mini	Clinical	120	120	.382	.382	4	4	13	13
	GPT-4o-mini	Narrative	120	120	.207	.207	4	4	13	13
	GLM-4.7	Clinical	120	80	.749	.715	6	6	21	23

Table S4, continued

Outcome	Model	Framing	Personas		$\phi(1)$		Minimum k		Recommended k	
			Full	Dec.	Full	Dec.	Full	Dec.	Full	Dec.
	GLM-4.7	Narrative	120	120	.850	.850	4	4	15	15
PHQ-8 $\geq$ 10	DeepSeek-V3	Clinical	120	111	.427	.424	12	12	45	45
	DeepSeek-V3	Narrative	120	120	.294	.294	9	9	34	34
	Gemini-3-Flash	Clinical	120	120	.745	.745	7	7	26	26
	Gemini-3-Flash	Narrative	120	120	.792	.792	6	6	21	21
	GPT-4o-mini	Clinical	120	120	.252	.252	16	16	63	63
	GPT-4o-mini	Narrative	120	120	.119	.119	18	18	71	71
	GLM-4.7	Clinical	120	80	.584	.550	10	11	39	41
	GLM-4.7	Narrative	120	120	.698	.698	8	8	29	29

Note. Dec. = December rows only (dec28_main, dec28_restored and all narrative rows). Minimum: the k for SE $\leq$ 0.25 reference SD (0.98 PHQ-8 points). Recommended: the k for SE $\leq$ 0.125 reference SD (0.49 points). Indicator rows: SE $\leq$ 0.10 and 0.05 in proportion.

Table S5

Level 1 on the Full Corpus and on December Rows Only

Model	Framing	Misfit ratio		Overfit ratio		Verdict, full	Verdict, Dec.
		Full	Dec.	Full	Dec.		
DeepSeek-V3	Both	0.05	0.05	0.00	0.00	compressed	compressed
DeepSeek-V3	Clinical	0.07	0.07	0.00	0.00	compressed	compressed
DeepSeek-V3	Narrative	0.03	0.03	0.00	0.00	compressed	compressed
Gemini-3-Flash	Both	0.00	0.00	0.17	0.17	compressed	compressed
Gemini-3-Flash	Clinical	0.00	0.00	0.16	0.16	compressed	compressed
Gemini-3-Flash	Narrative	0.00	0.00	0.18	0.18	compressed	compressed
GPT-4o-mini	Both	0.18	0.18	0.00	0.00	compressed	compressed
GPT-4o-mini	Clinical	0.20	0.20	0.00	0.00	compressed	compressed
GPT-4o-mini	Narrative	0.17	0.17	0.00	0.00	compressed	compressed
GLM-4.7	Both	0.02	0.01	0.77	0.78	misfit deficit	misfit deficit
GLM-4.7	Clinical	0.02	0.02	0.68	0.64	misfit deficit	misfit deficit
GLM-4.7	Narrative	0.01	0.01	0.87	0.87	misfit deficit	misfit deficit
Pooled	Both	0.06	0.06	0.24	0.22	compressed	compressed
Pooled	Clinical	0.07	0.08	0.21	0.16	compressed	compressed
Pooled	Narrative	0.05	0.05	0.26	0.26	compressed	compressed

Note. Reference: NHANES 2005–2018, weighted graded response model; all personas. Every verdict is a fail; the verdict column gives the failure reading, where compressed means a deficit in both tails.

Table S6

Level 2 on the Full Corpus and on December Rows Only

Estimand	Contrast	Model	Full corpus		December only		Changed
			Ratio	Verdict	Ratio	Verdict	
Standardised	Black minus White	DeepSeek-V3	−0.75	reference too imprecise	−0.82	reference too imprecise	

Table S6, continued

Estimand	Contrast	Model	Full corpus		December only		Changed
			Ratio	Verdict	Ratio	Verdict	
		Gemini-3-Flash	1.39	reference too imprecise	1.39	reference too imprecise	
		GPT-4o-mini	−0.76	reversed or missing	−0.76	reversed or missing	
		GLM-4.7	1.96	reference too imprecise	1.96	reference too imprecise	
		Pooled	0.46	reference too imprecise	0.44	reference too imprecise	
	Hispanic minus White	DeepSeek-V3	−0.95	no population gap	−0.95	no population gap	
		Gemini-3-Flash	0.04	no population gap	0.04	no population gap	
		GPT-4o-mini	−0.77	no population gap	−0.77	no population gap	
		GLM-4.7	2.42	no population gap	3.33	no population gap	
		Pooled	0.19	no population gap	0.41	no population gap	
	Asian minus White	DeepSeek-V3	0.74	attenuated or kept	0.74	attenuated or kept	
		Gemini-3-Flash	1.52	undetermined	1.52	undetermined	
		GPT-4o-mini	0.42	undetermined	0.42	undetermined	
		GLM-4.7	1.82	undetermined	1.82	undetermined	
		Pooled	1.13	undetermined	1.13	undetermined	
	Women minus Men	DeepSeek-V3	0.41	missing or attenuated	0.37	missing or attenuated	
		Gemini-3-Flash	1.17	undetermined	1.17	undetermined	
		GPT-4o-mini	0.23	missing or attenuated	0.23	missing or attenuated	
		GLM-4.7	2.06	steepened	2.00	steepened	
		Pooled	0.97	undetermined	0.94	attenuated or kept	
	Low minus High SES	DeepSeek-V3	1.91	steepened	1.94	steepened	
		Gemini-3-Flash	5.07	steepened	5.07	steepened	
		GPT-4o-mini	2.16	steepened	2.16	steepened	
		GLM-4.7	5.03	steepened	5.14	steepened	
		Pooled	3.54	steepened	3.58	steepened	
	Middle minus High SES	DeepSeek-V3	2.57	steepened	2.65	steepened	
		Gemini-3-Flash	1.31	reference too imprecise	1.31	reference too imprecise	
		GPT-4o-mini	2.25	steepened	2.25	steepened	
		GLM-4.7	2.21	steepened	2.25	steepened	
		Pooled	2.08	steepened	2.11	steepened	
	Low minus Middle SES	DeepSeek-V3	1.33	reference too imprecise	1.32	reference too imprecise	
		Gemini-3-Flash	8.38	steepened	8.38	steepened	
		GPT-4o-mini	2.08	steepened	2.08	steepened	
		GLM-4.7	7.51	steepened	7.68	steepened	
		Pooled	4.83	steepened	4.86	steepened	
Marginal	Black minus White	DeepSeek-V3	0.65	reference too imprecise	0.71	reference too imprecise	
		Gemini-3-Flash	−1.20	reference too imprecise	−1.20	reference too imprecise	
		GPT-4o-mini	0.66	reference too imprecise	0.66	reference too imprecise	
		GLM-4.7	−1.69	reversed	−1.69	reversed	
		Pooled	−0.40	reversed or missing	−0.38	reversed or missing	
	Hispanic minus White	DeepSeek-V3	0.72	reference too imprecise	0.72	reference too imprecise	
		Gemini-3-Flash	−0.03	reference too imprecise	−0.03	reference too imprecise	
		GPT-4o-mini	0.58	reference too imprecise	0.58	reference too imprecise	

Table S6, continued

Estimand	Contrast	Model	Full corpus		December only		Changed
			Ratio	Verdict	Ratio	Verdict	
	Asian minus White	GLM-4.7	−1.83	reversed or missing	−2.52	reversed or missing	
		Pooled	−0.14	reference too imprecise	−0.31	reference too imprecise	
		DeepSeek-V3	0.94	reference too imprecise	0.94	reference too imprecise	
		Gemini-3-Flash	1.92	reference too imprecise	1.92	reference too imprecise	
		GPT-4o-mini	0.53	reference too imprecise	0.53	reference too imprecise	
	Women minus Men	GLM-4.7	2.30	reference too imprecise	2.30	reference too imprecise	
		Pooled	1.42	reference too imprecise	1.42	reference too imprecise	
		DeepSeek-V3	0.37	missing or attenuated	0.34	missing or attenuated	
		Gemini-3-Flash	1.05	undetermined	1.05	undetermined	
		GPT-4o-mini	0.20	missing or attenuated	0.20	missing or attenuated	
	Low minus High SES	GLM-4.7	1.84	steepened	1.79	steepened	
		Pooled	0.87	attenuated or kept	0.85	attenuated or kept	
		DeepSeek-V3	1.23	kept or steepened	1.25	kept or steepened	
		Gemini-3-Flash	3.26	steepened	3.26	steepened	
		GPT-4o-mini	1.39	kept or steepened	1.39	kept or steepened	
	Middle minus High SES	GLM-4.7	3.23	steepened	3.30	steepened	
		Pooled	2.28	steepened	2.30	steepened	
		DeepSeek-V3	1.75	steepened	1.81	steepened	
		Gemini-3-Flash	0.89	attenuated or kept	0.89	attenuated or kept	
		GPT-4o-mini	1.53	kept or steepened	1.53	kept or steepened	
	Low minus Middle SES	GLM-4.7	1.50	kept or steepened	1.53	kept or steepened	
		Pooled	1.42	kept or steepened	1.44	kept or steepened	
		DeepSeek-V3	0.82	attenuated or kept	0.80	attenuated or kept	
		Gemini-3-Flash	5.13	steepened	5.13	steepened	
		GPT-4o-mini	1.27	kept or steepened	1.27	kept or steepened	
		GLM-4.7	4.60	steepened	4.70	steepened	
		Pooled	2.95	steepened	2.98	steepened	

Note. Rows of `12_verdicts.csv` with analysis `headline` and `dec_only`. Ratio = simulated gap over population gap. * = verdict differs.

Table S7

Level 3 on the Full Corpus and on December Rows Only

Model	Overall residual		Group residuals, range		Largest change	All groups pass at 2 points	
	Full	Dec.	Full	Dec.		Full	Dec.
DeepSeek-V3	4.58	4.55	4.06 to 5.13	4.06 to 5.10	0.25	No	No
Gemini-3-Flash	2.59	2.59	1.33 to 7.37	1.33 to 7.37	0.00	No	No
GPT-4o-mini	4.68	4.68	4.23 to 5.32	4.23 to 5.32	0.00	No	No
GLM-4.7	2.12	2.15	0.66 to 6.40	0.67 to 6.41	0.75	No	No
Pooled	3.49	3.50	2.59 to 6.05	2.61 to 6.07	0.15	No	No

Note. Specification S1, post-stratified estimand, framings combined, PHQ-8 mean in points. Largest change = largest absolute change in a group residual across the ten groups.

Table S8*Level 4 on the Full Corpus and on December Rows Only*

Model	Framing	λ_1/λ_2		Tucker ϕ		R4 at .08		Level 4 pass	
		Full	Dec.	Full	Dec.	Full	Dec.	Full	Dec.
DeepSeek-V3	Clinical	1.73	1.72	.349	.351	not interpretable	not interpretable	No	No
DeepSeek-V3	Narrative	1.43	1.43	.163	.163	not interpretable	not interpretable	No	No
Gemini-3-Flash	Clinical	8.69	8.69	.990	.990	fail	fail	No	No
Gemini-3-Flash	Narrative	9.33	9.33	.990	.990	fail	fail	No	No
GPT-4o-mini	Clinical	1.09	1.09	.364	.364	not interpretable	not interpretable	No	No
GPT-4o-mini	Narrative	1.33	1.33	.043	.043	not interpretable	not interpretable	No	No
GLM-4.7	Clinical	5.73	5.23	.996	.994	fail	fail	No	No
GLM-4.7	Narrative	8.12	8.12	.998	.998	fail	fail	No	No

Note. R1 needs $\lambda_1/\lambda_2 \geq 3$ and every loading $\geq .30$; R2 reads the lower 90% limit of Tucker's ϕ against .95.

S3.3 NHANES 2021–2023 Reference

The gate reads no population reference, so it has no 2021–2023 version. Level 1 was refitted with a graded response model on NHANES 2021–2023 (Table S9). Under that reference the level-1 verdicts that change, on both framings pooled, are: DeepSeek-V3, compressed to overfit deficit; GPT-4o-mini, compressed to overfit deficit. Level 3 was rerun against NHANES 2021–2023 and the 2017–2020 pre-pandemic file, each with its own design (Tables S10 and S11). Against 2021–2023 the per-model overall residuals move by -2.09 to -0.93 points (mean -1.16), and the per-model group residuals range from -0.33 to 6.27 points. Level 2 and level 4 were not run against NHANES 2021–2023.

Table S9*Level 1 Against NHANES 2005–2018 and NHANES 2021–2023*

Model	Framing	Misfit ratio		Overfit ratio		Verdict, 2005–18	Verdict, 2021–23
		2005–18	2021–23	2005–18	2021–23		
DeepSeek-V3	Both	0.05	0.28	0.00	0.00	compressed	overfit deficit
DeepSeek-V3	Clinical	0.07	0.34	0.00	0.00	compressed	overfit deficit
DeepSeek-V3	Narrative	0.03	0.22	0.00	0.00	compressed	overfit deficit
Gemini-3-Flash	Both	0.00	0.26	0.17	0.12	compressed	compressed
Gemini-3-Flash	Clinical	0.00	0.33	0.16	0.11	compressed	compressed
Gemini-3-Flash	Narrative	0.00	0.18	0.18	0.13	compressed	compressed
GPT-4o-mini	Both	0.18	0.99	0.00	0.00	compressed	overfit deficit

Table S9, continued

Model	Framing	Misfit ratio		Overfit ratio		Verdict, 2005–18	Verdict, 2021–23
		2005–18	2021–23	2005–18	2021–23		
GPT-4o-mini	Clinical	0.20	1.00	0.00	0.00	compressed	overfit deficit
GPT-4o-mini	Narrative	0.17	0.97	0.00	0.00	compressed	overfit deficit
GLM-4.7	Both	0.02	0.06	0.77	0.83	misfit deficit	misfit deficit
GLM-4.7	Clinical	0.02	0.07	0.68	0.74	misfit deficit	misfit deficit
GLM-4.7	Narrative	0.01	0.05	0.87	0.91	misfit deficit	misfit deficit
pooled	Both	0.06	0.40	0.24	0.24	compressed	overfit deficit
pooled	Clinical	0.07	0.44	0.21	0.21	compressed	overfit deficit
pooled	Narrative	0.05	0.36	0.26	0.26	compressed	overfit deficit

Note. Each reference has its own weighted graded response model. Every verdict is a fail; the verdict column gives the failure reading, where compressed means a deficit in both tails.

Table S10

Level 3 Overall Residual Under Three NHANES Reference Periods

Model	2005–2018 (S1)			2017–2020 (S6)			2021–2023 (S7)		
	Overall	Row	All	Overall	Row	All	Overall	Row	All
DeepSeek-V3	4.58	No	No	4.30	No	No	3.47	No	No
Gemini-3-Flash	2.59	No	No	2.22	No	No	1.36	Yes	No
GPT-4o-mini	4.68	No	No	4.41	No	No	3.57	No	No
GLM-4.7	2.12	No	No	1.74	Yes	No	0.91	Yes	No
Pooled	3.49	No	No	3.17	No	No	2.33	No	No

Note. Overall = post-stratified overall residual (simulated minus reference, PHQ-8 points), framings combined. Row = the overall row passes at 2 points. All = all ten group rows pass at 2 points.

Table S11

Level 3 Pooled Residual by Group Under Three Reference Periods

Group	Reference mean			Pooled residual		Shift
	2005–18	2017–20	2021–23	2005–18	2021–23	
Overall	2.89	3.10	3.96	3.49	2.33	–1.17
White	2.86	3.15	3.91	3.26	2.17	–1.09
Black	3.20	3.10	4.21	4.11	2.99	–1.12
Asian	2.16	2.08	3.10	3.34	2.16	–1.18
Hispanic	3.13	3.22	4.32	4.33	2.63	–1.70
Men	2.43	2.59	3.44	3.43	2.31	–1.12
Women	3.34	3.57	4.47	3.55	2.35	–1.21
Low	4.14	4.39	5.24	6.05	4.95	–1.10
Middle	3.07	3.43	4.33	3.16	1.91	–1.24
High	2.18	2.35	3.15	2.59	1.62	–0.98

Note. PHQ-8 points, post-stratified, pooled over models and framings. Shift = residual under 2021–2023 minus residual under 2005–2018.

S3.4 Marital-Status Codings

The primary reference counts married adults as married and widowed, divorced, separated and never-married adults as single, leaves cohabiting adults out, and counts 18- and 19-year-olds as single. Level 3 was rerun with cohabiting adults counted as married (S2) and as single (S2b) (Table S12). Level 2 was rerun with four other references: every marital status kept (all marital), single read as never married, 18- and 19-year-olds dropped, and marital status left out of the standardisation (Table S13). Across those four references, 25 per-model standardised verdict codes differ from the primary.

Table S12

Level 3 Under Three Codings of Marital Status

Model	S1: cohabiting dropped			S2: as married			S2b: as single		
	Overall	Max $ R $	Pass	Overall	Max $ R $	Pass	Overall	Max $ R $	Pass
DeepSeek-V3	4.58	5.13	No	4.55	5.08	No	4.59	5.13	No
Gemini-3-Flash	2.59	7.37	No	2.59	7.23	No	2.66	7.39	No
GPT-4o-mini	4.68	5.32	No	4.65	5.25	No	4.68	5.30	No
GLM-4.7	2.12	6.40	No	2.10	6.24	No	2.20	6.41	No
Pooled	3.49	6.05	No	3.47	5.94	No	3.53	6.05	No

Note. Post-stratified residual in PHQ-8 points, framings combined. Max $|R|$ = largest absolute group residual over ten groups. Pass = all ten group rows pass at 2 points.

Table S13

Level 2 Standardised Ratios Under Alternative Marital-Status References

Contrast	Model	Primary	All marital	Never married	Drop 18–19	No marital
Black minus White	DeepSeek-V3	−0.75 X	−0.71 X	−1.14 N	−0.75 X	−1.01 X
	Gemini-3-Flash	1.39 X	1.31 X	2.10 N	1.38 X	1.86 X
	GPT-4o-mini	−0.76 r/m	−0.72 r/m	−1.15 N	−0.76 r/m	−1.02 X
	GLM-4.7	1.96 X	1.85 X	2.97 N	1.96 X	2.63 X
	Pooled	0.46 X	0.43 X	0.70 N	0.46 X	0.62 X
Hispanic minus White	DeepSeek-V3	−0.95 N	−0.62 X	−1.15 N	−0.91 N	−0.60 X
	Gemini-3-Flash	0.04 N	0.03 X	0.05 N	0.04 N	0.02 X
	GPT-4o-mini	−0.77 N	−0.50 r/m	−0.93 N	−0.73 N	−0.49 r/m
	GLM-4.7	2.42 N	1.59 X	2.93 N	2.31 N	1.54 X
	Pooled	0.19 N	0.12 X	0.22 N	0.18 N	0.12 X
Asian minus White	DeepSeek-V3	0.74 a/k	0.72 a/k	0.81 X	0.72 X	0.68 a/k

Table S13, continued

Contrast	Model	Primary	All marital	Never married	Drop 18–19	No marital
	Gemini-3-Flash	1.52 u	1.47 u	1.65 X	1.48 X	1.38 u
	GPT-4o-mini	0.42 u	0.41 u	0.46 X	0.41 X	0.38 m/a
	GLM-4.7	1.82 u	1.76 u	1.98 X	1.77 X	1.66 u
	Pooled	1.13 u	1.09 u	1.22 X	1.10 X	1.03 u
Women minus Men	DeepSeek-V3	0.41 m/a	0.40 m/a	0.45 m/a	0.41 m/a	0.39 m/a
	Gemini-3-Flash	1.17 u	1.13 u	1.28 k/s	1.17 u	1.11 u
	GPT-4o-mini	0.23 m/a	0.22 m/a	0.25 m/a	0.23 m/a	0.22 m/a
	GLM-4.7	2.06 s	1.98 s	2.25 s	2.06 s	1.96 s
	Pooled	0.97 u	0.93 a/k	1.06 k/s	0.97 u	0.92 a/k
Low minus High SES	DeepSeek-V3	1.91 s	1.87 s	2.12 s	1.83 s	1.66 s
	Gemini-3-Flash	5.07 s	4.96 s	5.62 s	4.85 s	4.42 s
	GPT-4o-mini	2.16 s	2.11 s	2.39 s	2.06 s	1.88 s
	GLM-4.7	5.03 s	4.92 s	5.58 s	4.81 s	4.39 s
	Pooled	3.54 s	3.47 s	3.93 s	3.39 s	3.09 s
Middle minus High SES	DeepSeek-V3	2.57 s	2.55 s	3.13 s	2.50 s	2.30 s
	Gemini-3-Flash	1.31 X	1.30 X	1.60 X	1.27 X	1.17 X
	GPT-4o-mini	2.25 s	2.23 s	2.75 s	2.20 s	2.02 s
	GLM-4.7	2.21 s	2.19 s	2.70 s	2.15 s	1.98 X
	Pooled	2.08 s	2.07 s	2.54 s	2.03 s	1.87 s
Low minus Middle SES	DeepSeek-V3	1.33 X	1.29 X	1.37 X	1.25 X	1.14 u
	Gemini-3-Flash	8.38 s	8.10 s	8.59 s	7.87 s	7.13 s
	GPT-4o-mini	2.08 s	2.01 s	2.13 s	1.95 s	1.77 s
	GLM-4.7	7.51 s	7.26 s	7.70 s	7.06 s	6.39 s
	Pooled	4.83 s	4.67 s	4.95 s	4.53 s	4.11 s

Note. Each cell: ratio and verdict code. Codes: **s** steepened, **k/s** kept or steepened, **a/k** attenuated or kept, **m/a** missing or attenuated, **r** reversed, **r/m** reversed or missing, **u** undetermined, **X** reference too imprecise, **N** no population gap. Rows of `12_verdicts.csv` by analysis.

S3.5 Income-Anchor Sensitivity

The persona labels state dollar incomes and insurance, and the primary reference bands the NHANES ratio of family income to the poverty line. The sensitivity mapping uses household income and insurance type on the 2007–2018 cycles, where NHANES top-codes household income. Table S14 gives the level-3 income-group residuals under each mapping, Table S15 the marginal residuals, and Table S16 the level-2 income contrasts.

Table S14*Level 3 Income-Group Residuals Under Four Income Mappings*

Model	S1: poverty ratio		S3: PIR 2007–18		S4: persona string		S5: top-coded	
	Low, Mid, High	Pass	Low, Mid, High	Pass	Low, Mid, High	Pass	Low, Mid, High	Pass
DeepSeek-V3	5.13, 4.92, 4.06	No	5.08, 4.84, 4.02	No	3.98, 5.43, 4.21	No	3.98, 5.43, 4.27	No
Gemini-3-Flash	7.37, 1.33, 1.42	No	7.32, 1.25, 1.37	No	6.32, 1.71, 1.59	No	6.32, 1.71, 1.65	No
GPT-4o-mini	5.32, 4.88, 4.23	No	5.26, 4.79, 4.19	No	4.13, 5.32, 4.43	No	4.13, 5.32, 4.50	No
GLM-4.7	6.40, 1.50, 0.66	No	6.36, 1.43, 0.61	No	5.45, 1.77, 0.79	No	5.45, 1.77, 0.85	No
Pooled	6.05, 3.16, 2.59	No	6.01, 3.08, 2.55	No	4.97, 3.56, 2.75	No	4.97, 3.56, 2.82	No

Note. Post-stratified residuals in PHQ-8 points, framings combined. S3 uses the poverty ratio on 2007–2018 only; S4 maps the persona's dollar income and insurance onto household income and insurance type; S5 is S4 with the high band top-coded. Pass = all ten group rows pass at 2 points.

Table S15*Marginal Income-Group Residuals Under Two Income Mappings*

Model	Low		Middle		High	
	S3	S4	S3	S4	S3	S4
DeepSeek-V3	5.08	3.98	4.84	5.43	4.02	4.21
Gemini-3-Flash	7.32	6.32	1.25	1.71	1.37	1.59
GPT-4o-mini	5.26	4.13	4.79	5.32	4.19	4.43
GLM-4.7	6.36	5.45	1.43	1.77	0.61	0.79

Note. Marginal residual (simulated group mean minus reference group mean), PHQ-8 points, from 80d_income_anchor_permodel.csv.

Table S16*Level 2 Income Contrasts Under the Alternative Income Mapping*

Estimand	Contrast	Model	Poverty ratio (primary)		Persona string	
			Ratio	Verdict	Ratio	Verdict
Standardised	Low minus High SES	DeepSeek-V3	1.91	steepened	1.04	kept or steepened
		Gemini-3-Flash	5.07	steepened	2.76	steepened
		GPT-4o-mini	2.16	steepened	1.18	kept or steepened
		GLM-4.7	5.03	steepened	2.74	steepened
		Pooled	3.54	steepened	1.93	steepened
	Middle minus High SES	DeepSeek-V3	2.57	steepened	3.29	steepened
		Gemini-3-Flash	1.31	reference too imprecise	1.68	reference too imprecise
		GPT-4o-mini	2.25	steepened	2.89	steepened
		GLM-4.7	2.21	steepened	2.83	steepened
		Pooled	2.08	steepened	2.67	steepened
	Low minus Middle SES	DeepSeek-V3	1.33	reference too imprecise	0.48	attenuated
		Gemini-3-Flash	8.38	steepened	3.03	steepened

Table S16, continued

Estimand	Contrast	Model	Poverty ratio (primary)		Persona string	
			Ratio	Verdict	Ratio	Verdict
Marginal	Low minus High SES	GPT-4o-mini	2.08	steepened	0.75	attenuated or kept
		GLM-4.7	7.51	steepened	2.72	steepened
		Pooled	4.83	steepened	1.75	steepened
	Middle minus High SES	DeepSeek-V3	1.23	kept or steepened	0.70	attenuated or kept
		Gemini-3-Flash	3.26	steepened	1.86	steepened
		GPT-4o-mini	1.39	kept or steepened	0.79	attenuated or kept
		GLM-4.7	3.23	steepened	1.85	steepened
		Pooled	2.28	steepened	1.30	kept or steepened
		DeepSeek-V3	1.75	steepened	3.07	steepened
		Gemini-3-Flash	0.89	attenuated or kept	1.56	reference too imprecise
		GPT-4o-mini	1.53	kept or steepened	2.69	steepened
		GLM-4.7	1.50	kept or steepened	2.64	steepened
		Pooled	1.42	kept or steepened	2.49	steepened
	Low minus Middle SES	DeepSeek-V3	0.82	attenuated or kept	0.30	missing or attenuated
		Gemini-3-Flash	5.13	steepened	1.91	steepened
		GPT-4o-mini	1.27	kept or steepened	0.47	attenuated
		GLM-4.7	4.60	steepened	1.71	steepened
		Pooled	2.95	steepened	1.10	kept or steepened

Note. Rows of `12_verdicts.csv` with analysis `headline` and `ref_income_alt`. The alternative mapping was run for the marginal estimand on the income contrasts only.

S4. Full Per-Model Results

S4.1 Gate

Tables S17 and S18 give every gate quantity with its persona-bootstrap interval, for the PHQ-8 total and for the indicator of a total of 10 or more.

Table S17

Gate: Variance Components and Dependability With Persona-Bootstrap Intervals

Outcome	Model	Framing	σ_p^2	σ_r^2	$\phi(1)$	$\phi(30)$
PHQ-8 total	DeepSeek-V3	Clinical	2.14 [1.58, 2.71]	2.03 [1.78, 2.28]	.513 [.449, .563]	.969 [.961, .975]
	DeepSeek-V3	Narrative	1.61 [1.16, 2.09]	1.57 [1.39, 1.78]	.505 [.439, .561]	.968 [.959, .975]
	Gemini-3-Flash	Clinical	18.74 [15.78, 21.32]	2.17 [1.90, 2.42]	.896 [.870, .916]	.996 [.995, .997]
	Gemini-3-Flash	Narrative	20.69 [17.78, 23.42]	2.27 [2.02, 2.54]	.901 [.877, .919]	.996 [.995, .997]
	GPT-4o-mini	Clinical	1.89 [1.51, 2.23]	3.06 [2.85, 3.27]	.382 [.332, .419]	.949 [.937, .956]
	GPT-4o-mini	Narrative	0.81 [0.62, 0.98]	3.08 [2.88, 3.32]	.207 [.168, .243]	.887 [.858, .906]

Table S17, continued

Outcome	Model	Framing	σ_p^2	σ_r^2	$\phi(1)$	$\phi(30)$
PHQ-8 ≥ 10	GLM-4.7	Clinical	15.04 [12.10, 17.66]	5.04 [4.56, 5.53]	.749 [.697, .787]	.989 [.986, .991]
	GLM-4.7	Narrative	20.16 [16.68, 23.27]	3.55 [3.25, 3.88]	.850 [.814, .874]	.994 [.992, .995]
	DeepSeek-V3	Clinical	0.08 [0.06, 0.10]	0.11 [0.10, 0.13]	.426 [.349, .499]	.957 [.941, .968]
	DeepSeek-V3	Narrative	0.04 [0.02, 0.05]	0.08 [0.07, 0.10]	.294 [.224, .360]	.926 [.897, .944]
	Gemini-3-Flash	Clinical	0.19 [0.17, 0.20]	0.06 [0.05, 0.08]	.745 [.677, .802]	.989 [.984, .992]
	Gemini-3-Flash	Narrative	0.20 [0.18, 0.21]	0.05 [0.04, 0.07]	.792 [.730, .845]	.991 [.988, .994]
	GPT-4o-mini	Clinical	0.05 [0.04, 0.06]	0.16 [0.14, 0.17]	.252 [.210, .291]	.910 [.889, .925]
	GPT-4o-mini	Narrative	0.02 [0.02, 0.03]	0.18 [0.16, 0.19]	.119 [.092, .143]	.801 [.752, .834]
	GLM-4.7	Clinical	0.14 [0.12, 0.15]	0.10 [0.08, 0.12]	.584 [.507, .652]	.977 [.969, .983]
	GLM-4.7	Narrative	0.16 [0.14, 0.18]	0.07 [0.05, 0.09]	.698 [.620, .765]	.986 [.980, .990]

Note. Brackets: 95% persona-bootstrap intervals. σ_p^2 = variance between personas, σ_r^2 = variance between draws of one persona (squared PHQ-8 points for the total, squared proportion for the indicator). ϕ from the nested design.

Table S18

Gate: Draws Needed per Persona With Persona-Bootstrap Intervals

Outcome	Model	Framing	Separation		Gate rule	
			$\phi \geq .80$	$\phi \geq .90$	Min.	Rec.
PHQ-8 total	DeepSeek-V3	Clinical	4 [4, 5]	9 [7, 12]	3 [2, 3]	9 [8, 10]
	DeepSeek-V3	Narrative	4 [4, 6]	9 [8, 12]	2 [2, 2]	7 [6, 8]
	Gemini-3-Flash	Clinical	1 [1, 1]	2 [1, 2]	3 [2, 3]	9 [8, 10]
	Gemini-3-Flash	Narrative	1 [1, 1]	1 [1, 2]	3 [3, 3]	10 [9, 11]
	GPT-4o-mini	Clinical	7 [6, 9]	15 [13, 19]	4 [3, 4]	13 [12, 14]
	GPT-4o-mini	Narrative	16 [13, 20]	35 [29, 45]	4 [3, 4]	13 [12, 14]
	GLM-4.7	Clinical	2 [2, 2]	4 [3, 4]	6 [5, 6]	21 [19, 23]
	GLM-4.7	Narrative	1 [1, 1]	2 [2, 3]	4 [4, 5]	15 [14, 17]
PHQ-8 ≥ 10	DeepSeek-V3	Clinical	6 [5, 8]	13 [10, 17]	12 [10, 13]	45 [39, 52]
	DeepSeek-V3	Narrative	10 [8, 14]	22 [16, 32]	9 [7, 11]	34 [28, 41]
	Gemini-3-Flash	Clinical	2 [1, 2]	4 [3, 5]	7 [5, 9]	26 [20, 33]
	Gemini-3-Flash	Narrative	2 [1, 2]	3 [2, 4]	6 [4, 7]	21 [16, 28]
	GPT-4o-mini	Clinical	12 [10, 16]	27 [22, 34]	16 [14, 17]	63 [56, 68]
	GPT-4o-mini	Narrative	30 [24, 40]	67 [54, 90]	18 [17, 20]	71 [66, 77]
	GLM-4.7	Clinical	3 [3, 4]	7 [5, 9]	10 [9, 12]	39 [33, 47]
	GLM-4.7	Narrative	2 [2, 3]	4 [3, 6]	8 [6, 9]	29 [22, 36]

Note. Brackets: 95% persona-bootstrap intervals for k . Min. = k for SE ≤ 0.25 reference SD (0.98 PHQ-8 points); Rec. = k for SE ≤ 0.125 reference SD (0.49 points); reference SD 3.935 (NHANES 2005–2018). Indicator rows: SE ≤ 0.10 and 0.05 in proportion. Separation columns: k for persona separation $\phi(k)$, required only for uses that compare personas.

S4.1a Level 1

Table S20 gives the person-fit tails for each model and framing.

S4.2 Level 2

Table S21 gives the standardised estimand, which is primary, and Table S22 the marginal estimand. Each table carries the verdict under both stops and, as a sensitivity

Table S19*Gate: Precision and Separation of a Persona Mean of k Draws, PHQ-8 Total*

Model	Framing	Variance		SE(30)	Separation		k needed		Band flip (%)
		σ_p^2	σ_r^2		$\phi(1)$	$\phi(30)$	Minimum	Recommended	
DeepSeek-V3	Clinical	2.14	2.03	0.26	.51	.97	3	9	27.0
	Narrative	1.61	1.57	0.23	.51	.97	2	7	22.4
Gemini-3-Flash	Clinical	18.74	2.17	0.27	.90	>.99	3	9	31.6
	Narrative	20.69	2.27	0.27	.90	>.99	3	10	28.7
GPT-4o-mini	Clinical	1.89	3.06	0.32	.38	.95	4	13	37.3
	Narrative	0.81	3.09	0.32	.21	.89	4	13	39.9
GLM-4.7	Clinical	15.04	5.04	0.41	.75	.99	6	21	44.8
	Narrative	20.16	3.55	0.34	.85	.99	4	15	37.6

Note. σ_p^2 = variance between personas; σ_r^2 = variance between draws of one persona, in squared PHQ-8 points. SE(30) = standard error of a 30-draw persona mean, in points. $\phi(k)$ = dependability of a mean of k draws (Equation 2), the separation of personas, required only for uses that compare individual personas. Minimum: the smallest k with $SE(k) \leq 0.25$ NHANES standard deviations (0.98 points). Recommended: the smallest k with $SE(k) \leq 0.125$ standard deviations (0.49 points). Band flip = probability that two draws of one persona fall in different PHQ-8 severity categories, a description of single-draw labels. Persona-bootstrap intervals for every entry are in Supplement S4.

Table S20*Level 1: Person Fit of Single Draws Against NHANES Adults at the Same Total*

Model	Framing	Below NHANES 5th pct.		Above NHANES 95th pct.		Mid-PIT	Verdict
		%	Ratio 90% CI	%	Ratio 90% CI		
DeepSeek-V3	Clinical	0.3	0.01–0.21	0.0	0.00–0.01	.36	Compressed
	Narrative	0.1	0.00–0.11	0.0	0.00–0.01	.37	Compressed
Gemini-3-Flash	Clinical	0.0	0.00–0.00	0.8	0.11–0.39	.66	Compressed
	Narrative	0.0	0.00–0.00	0.9	0.11–0.42	.65	Compressed
GPT-4o-mini	Clinical	1.0	0.09–0.37	0.0	0.00–0.03	.38	Compressed
	Narrative	0.8	0.09–0.30	0.0	0.00–0.03	.38	Compressed
GLM-4.7	Clinical	0.1	0.01–0.04	3.4	0.55–0.92	.65	Misfit deficit
	Narrative	0.1	0.00–0.04	4.3	0.74–1.15	.65	Misfit deficit
NHANES		5.0		5.0		.50	

Note. Each draw is placed in the weighted NHANES 2005–2018 distribution of the person-fit statistic l_z^* among adults with the same PHQ-8 total. % = share of a model’s draws below the within-total 5th percentile (misfit tail) or above the 95th (overfit tail); NHANES places 5% in each by construction. Ratio = share divided by .05, with a 90% interval from a persona-clustered bootstrap drawn jointly with a Rao–Wu bootstrap of NHANES. A model passes when both ratio intervals lie inside 0.67 to 1.5 ($\tau = 1.5$). Mid-PIT = mean probability integral transform, .50 for a population like NHANES. Compressed = a deficit in both tails.

column, the verdict under the population-only stop.

Figure S1 draws every per-model and pooled ratio of the article's Figure 2 with its 98.6% Fieller interval.

Table S21

Level 2 per Model, Standardised Estimand

Contrast	Model	Pop. gap	Sim. gap	Ratio	98.6% interval	Verdict	Population-only stop
Black minus White	DeepSeek-V3	−0.28	0.21	−0.75	[−3.74, 0.36]	reference too imprecise	reference too imprecise
	Gemini-3-Flash		−0.39	1.39	[−0.61, 6.84]	reference too imprecise	reference too imprecise
	GPT-4o-mini		0.22	−0.76	[−3.63, 0.19]	reversed or missing	reference too imprecise
	GLM-4.7		−0.56	1.96	[0.29, 8.75]	reference too imprecise	reference too imprecise
	Pooled		−0.13	0.46	[−0.22, 2.34]	reference too imprecise	reference too imprecise
Hispanic minus White	DeepSeek-V3	−0.18	0.17	−0.95	unbounded	no population gap	no population gap
	Gemini-3-Flash		−0.01	0.04	unbounded	no population gap	no population gap
	GPT-4o-mini		0.14	−0.77	unbounded	no population gap	no population gap
	GLM-4.7		−0.44	2.42	unbounded	no population gap	no population gap
	Pooled		−0.03	0.19	unbounded	no population gap	no population gap
Asian minus White	DeepSeek-V3	−0.95	−0.71	0.74	[0.36, 1.20]	attenuated or kept	attenuated or kept
	Gemini-3-Flash		−1.45	1.52	[−0.39, 3.54]	undetermined	undetermined
	GPT-4o-mini		−0.40	0.42	[0.09, 0.79]	undetermined	undetermined
	GLM-4.7		−1.74	1.82	[0.55, 3.25]	undetermined	undetermined
	Pooled		−1.07	1.13	[0.57, 1.80]	undetermined	undetermined
Women minus Men	DeepSeek-V3	0.84	0.35	0.41	[0.23, 0.62]	missing or attenuated	missing or attenuated
	Gemini-3-Flash		0.99	1.17	[0.73, 1.70]	undetermined	undetermined
	GPT-4o-mini		0.19	0.23	[0.00, 0.47]	missing or attenuated	missing or attenuated
	GLM-4.7		1.73	2.06	[1.50, 2.76]	steepened	steepened
	Pooled		0.81	0.97	[0.75, 1.26]	undetermined	undetermined
Low minus High SES	DeepSeek-V3	1.30	2.48	1.91	[1.55, 2.39]	steepened	steepened
	Gemini-3-Flash		6.58	5.07	[3.85, 6.60]	steepened	steepened
	GPT-4o-mini		2.80	2.16	[1.78, 2.68]	steepened	steepened
	GLM-4.7		6.53	5.03	[4.05, 6.33]	steepened	steepened
	Pooled		4.60	3.54	[2.94, 4.38]	steepened	steepened
Middle minus High SES	DeepSeek-V3	0.61	1.56	2.57	[1.82, 3.94]	steepened	reference too imprecise
	Gemini-3-Flash		0.79	1.31	[0.86, 2.06]	reference too imprecise	reference too imprecise
	GPT-4o-mini		1.37	2.25	[1.57, 3.49]	steepened	reference too imprecise
	GLM-4.7		1.34	2.21	[1.32, 3.62]	steepened	reference too imprecise
	Pooled		1.26	2.08	[1.52, 3.17]	steepened	reference too imprecise
Low minus Middle SES	DeepSeek-V3	0.69	0.92	1.33	[0.83, 2.09]	reference too imprecise	reference too imprecise
	Gemini-3-Flash		5.79	8.38	[5.88, 12.59]	steepened	reference too imprecise

Table S21, continued

Contrast	Model	Pop. gap	Sim. gap	Ratio	98.6% interval	Verdict	Population-only stop
	GPT-4o-mini		1.44	2.08	[1.55, 3.04]	steepened	reference too imprecise
	GLM-4.7		5.19	7.51	[5.62, 10.96]	steepened	reference too imprecise
	Pooled		3.33	4.83	[3.62, 7.03]	steepened	reference too imprecise

Note. Gaps in PHQ-8 points. Ratio = simulated gap over population gap, with its Fieller interval at the Bonferroni level for 7 contrasts (one-sided $\alpha = .0071$). Verdict: the primary rule, whose reference-imprecision stop fires when the observed simulated gap, treated as measured without error, would still give a ratio interval spanning three or more regions. Population-only stop: the verdict when that stop instead fires whenever the population gap's relative half-width exceeds 0.25, before the simulation is read (sensitivity). Rows of `12_verdicts.csv` with analysis **headline**.

Table S22

Level 2 per Model, Marginal Estimand

Contrast	Model	Pop. gap	Sim. gap	Ratio	98.6% interval	Verdict	Population-only stop
Black minus White	DeepSeek-V3	0.33	0.21	0.65	[−0.31, 2.32]	reference too imprecise	reference too imprecise
	Gemini-3-Flash		−0.39	−1.20	[−4.24, 0.52]	reference too imprecise	reference too imprecise
	GPT-4o-mini		0.22	0.66	[−0.16, 2.21]	reference too imprecise	reference too imprecise
	GLM-4.7		−0.56	−1.69	[−5.15, −0.25]	reversed	reference too imprecise
	Pooled		−0.13	−0.40	[−1.46, 0.18]	reversed or missing	reference too imprecise
Hispanic minus White	DeepSeek-V3	0.24	0.17	0.72	[−0.79, 3.85]	reference too imprecise	reference too imprecise
	Gemini-3-Flash		−0.01	−0.03	[−1.54, 1.42]	reference too imprecise	reference too imprecise
	GPT-4o-mini		0.14	0.58	[−0.13, 2.65]	reference too imprecise	reference too imprecise
	GLM-4.7		−0.44	−1.83	[−8.15, 0.15]	reversed or missing	reference too imprecise
	Pooled		−0.03	−0.14	[−1.24, 0.62]	reference too imprecise	reference too imprecise
Asian minus White	DeepSeek-V3	−0.76	−0.71	0.94	[0.45, 1.53]	reference too imprecise	reference too imprecise
	Gemini-3-Flash		−1.45	1.92	[−0.50, 4.49]	reference too imprecise	reference too imprecise
	GPT-4o-mini		−0.40	0.53	[0.12, 1.00]	reference too imprecise	reference too imprecise
	GLM-4.7		−1.74	2.30	[0.69, 4.14]	reference too imprecise	reference too imprecise
	Pooled		−1.07	1.42	[0.72, 2.31]	reference too imprecise	reference too imprecise
Women minus Men	DeepSeek-V3	0.94	0.35	0.37	[0.20, 0.54]	missing or attenuated	missing or attenuated
	Gemini-3-Flash		0.99	1.05	[0.66, 1.47]	undetermined	undetermined
	GPT-4o-mini		0.19	0.20	[0.00, 0.41]	missing or attenuated	missing or attenuated
	GLM-4.7		1.73	1.84	[1.38, 2.36]	steepened	steepened
	Pooled		0.81	0.87	[0.69, 1.07]	attenuated or kept	attenuated or kept
Low minus High SES	DeepSeek-V3	2.02	2.48	1.23	[1.04, 1.44]	kept or steepened	kept or steepened
	Gemini-3-Flash		6.58	3.26	[2.54, 4.03]	steepened	steepened

Table S22, continued

Contrast	Model	Pop. gap	Sim. gap	Ratio	98.6% interval	Verdict	Population-only stop
Middle minus High SES	GPT-4o-mini		2.80	1.39	[1.21, 1.59]	kept or steepened	kept or steepened
	GLM-4.7		6.53	3.23	[2.71, 3.81]	steepened	steepened
	Pooled		4.60	2.28	[2.00, 2.60]	steepened	steepened
	DeepSeek-V3	0.89	1.56	1.75	[1.37, 2.21]	steepened	steepened
	Gemini-3-Flash		0.79	0.89	[0.63, 1.19]	attenuated or kept	attenuated or kept
	GPT-4o-mini		1.37	1.53	[1.17, 1.98]	kept or steepened	kept or steepened
	GLM-4.7		1.34	1.50	[0.94, 2.13]	kept or steepened	kept or steepened
	Pooled		1.26	1.42	[1.17, 1.75]	kept or steepened	kept or steepened
	DeepSeek-V3	1.13	0.92	0.82	[0.53, 1.16]	attenuated or kept	attenuated or kept
	Gemini-3-Flash		5.79	5.13	[3.81, 6.83]	steepened	steepened
Low minus Middle SES	GPT-4o-mini		1.44	1.27	[1.02, 1.62]	kept or steepened	kept or steepened
	GLM-4.7		5.19	4.60	[3.72, 5.83]	steepened	steepened
	Pooled		3.33	2.95	[2.40, 3.74]	steepened	steepened

Note. Gaps in PHQ-8 points. Ratio = simulated gap over population gap, with its Fieller interval at the Bonferroni level for 7 contrasts (one-sided $\alpha = .0071$). Verdict: the primary rule, whose reference-imprecision stop fires when the observed simulated gap, treated as measured without error, would still give a ratio interval spanning three or more regions. Population-only stop: the verdict when that stop instead fires whenever the population gap's relative half-width exceeds 0.25, before the simulation is read (sensitivity). Rows of `12_verdicts.csv` with analysis `headline`.

S4.3 Level 3

Table S23 gives every group row for the mean, and Table S24 every group row for the share scoring 10 or more.

Table S23

Level 3 per Model and Group: Mean PHQ-8

Model	Group	Sim.	Ref.	Residual	90% CI	Passes at tolerance			
						0.2 SD	1 pt	2 pt	5 pt (bench.)
DeepSeek-V3	Overall	7.48	2.89	4.58	[4.49, 4.67]	No	No	No	Yes
	White	7.28	2.86	4.42	[4.31, 4.54]	No	No	No	Yes
	Black	8.26	3.20	5.05	[4.91, 5.20]	No	No	No	No
	Asian	6.63	2.16	4.47	[4.32, 4.61]	No	No	No	Yes
	Hispanic	8.22	3.13	5.09	[4.96, 5.22]	No	No	No	No
	Men	7.23	2.43	4.80	[4.69, 4.92]	No	No	No	Yes
	Women	7.71	3.34	4.38	[4.26, 4.49]	No	No	No	Yes
	Low	9.27	4.14	5.13	[4.96, 5.30]	No	No	No	No
	Middle	7.99	3.07	4.92	[4.82, 5.02]	No	No	No	No
	High	6.24	2.18	4.06	[3.91, 4.20]	No	No	No	Yes
Gemini-3-Flash	Overall	5.48	2.89	2.59	[2.49, 2.69]	No	No	No	Yes
	White	5.09	2.86	2.23	[2.12, 2.35]	No	No	No	Yes
	Black	6.55	3.20	3.34	[3.16, 3.53]	No	No	No	Yes
	Asian	4.58	2.16	2.42	[2.26, 2.58]	No	No	No	Yes
	Hispanic	7.09	3.13	3.97	[3.77, 4.16]	No	No	No	Yes
	Men	4.91	2.43	2.48	[2.37, 2.59]	No	No	No	Yes
	Women	6.03	3.34	2.69	[2.56, 2.82]	No	No	No	Yes
	Low	11.51	4.14	7.37	[7.20, 7.53]	No	No	No	No

Table S23, continued

Model	Group	Sim.	Ref.	Residual	90% CI	Passes at tolerance			
						0.2 SD	1 pt	2 pt	5 pt (bench.)
GPT-4o-mini	Middle	4.40	3.07	1.33	[1.22, 1.44]	No	No	Yes	Yes
	High	3.60	2.18	1.42	[1.32, 1.52]	No	No	Yes	Yes
	Overall	7.57	2.89	4.68	[4.57, 4.79]	No	No	No	Yes
	White	7.37	2.86	4.51	[4.36, 4.66]	No	No	No	Yes
	Black	8.27	3.20	5.07	[4.90, 5.24]	No	No	No	No
	Asian	7.15	2.16	4.98	[4.81, 5.15]	No	No	No	No
	Hispanic	8.30	3.13	5.17	[5.02, 5.33]	No	No	No	No
	Men	7.34	2.43	4.91	[4.77, 5.05]	No	No	No	No
	Women	7.80	3.34	4.46	[4.30, 4.62]	No	No	No	Yes
	Low	9.46	4.14	5.32	[5.13, 5.51]	No	No	No	No
GLM-4.7	Middle	7.95	3.07	4.88	[4.73, 5.02]	No	No	No	No
	High	6.41	2.18	4.23	[4.04, 4.42]	No	No	No	Yes
	Overall	5.02	2.89	2.12	[2.00, 2.24]	No	No	No	Yes
	White	4.72	2.86	1.86	[1.71, 2.01]	No	No	No	Yes
	Black	6.16	3.20	2.96	[2.75, 3.17]	No	No	No	Yes
	Asian	3.66	2.16	1.50	[1.30, 1.69]	No	No	Yes	Yes
	Hispanic	6.20	3.13	3.08	[2.89, 3.27]	No	No	No	Yes
	Men	3.94	2.43	1.51	[1.39, 1.63]	No	No	Yes	Yes
	Women	6.03	3.34	2.69	[2.52, 2.86]	No	No	No	Yes
	Low	10.54	4.14	6.40	[6.21, 6.59]	No	No	No	No
Pooled	Middle	4.57	3.07	1.50	[1.34, 1.66]	No	No	Yes	Yes
	High	2.84	2.18	0.66	[0.51, 0.81]	No	Yes	Yes	Yes
	Overall	6.39	2.89	3.49	[3.42, 3.57]	No	No	No	Yes
	White	6.12	2.86	3.26	[3.16, 3.36]	No	No	No	Yes
	Black	7.31	3.20	4.11	[3.96, 4.25]	No	No	No	Yes
	Asian	5.51	2.16	3.34	[3.20, 3.48]	No	No	No	Yes
	Hispanic	7.45	3.13	4.33	[4.19, 4.47]	No	No	No	Yes
	Men	5.85	2.43	3.43	[3.34, 3.52]	No	No	No	Yes
	Women	6.89	3.34	3.55	[3.45, 3.66]	No	No	No	Yes
	Low	10.19	4.14	6.05	[5.91, 6.19]	No	No	No	No
	Middle	6.23	3.07	3.16	[3.06, 3.25]	No	No	No	Yes
	High	4.77	2.18	2.59	[2.50, 2.69]	No	No	No	Yes

Note. Post-stratified estimand, specification S1, framings combined. Residual = simulated minus reference, PHQ-8 points. 0.2 SD = 0.79 points. A row passes when its 90% interval lies inside the tolerance (two one-sided tests). The 5-point column is a benchmark only.

Table S24

Level 3 per Model and Group: Share Scoring 10 or More

Model	Group	Sim. %	Ref. %	Diff. (pp)	90% CI	Ratio	90% CI	Passes		
								2.5 pp	5 pp	0.80–1.25
DeepSeek-V3	Overall	7.4	7.0	0.3	[−1.0, 1.7]	1.05	[0.85, 1.25]	Yes	Yes	Yes
	White	6.1	6.8	−0.8	[−2.6, 1.1]	0.89	[0.62, 1.17]	No	Yes	No
	Black	14.0	8.8	5.1	[2.7, 7.6]	1.58	[1.30, 1.88]	No	No	No
	Asian	1.8	3.4	−1.6	[−3.4, 0.1]	0.53	[0.04, 1.05]	No	Yes	No
	Hispanic	11.2	8.1	3.1	[0.9, 5.3]	1.38	[1.10, 1.67]	No	No	No

Table S24, continued

Model	Group	Sim. %	Ref. %	Diff. (pp)	90% CI	Ratio	90% CI	Passes		
								2.5 pp	5 pp	0.80–1.25
	Men	4.8	5.2	−0.4	[−2.3, 1.6]	0.93	[0.56, 1.32]	Yes	Yes	No
	Women	9.7	8.7	1.0	[−0.9, 2.9]	1.11	[0.89, 1.34]	No	Yes	No
	Low	32.7	13.6	19.1	[16.0, 22.1]	2.40	[2.15, 2.68]	No	No	No
	Middle	1.4	7.3	−6.0	[−7.9, −4.0]	0.19	[0.00, 0.44]	No	No	No
	High	0.5	3.7	−3.2	[−5.5, −0.9]	0.14	[0.00, 0.75]	No	No	No
Gemini-3-Flash	Overall	16.4	7.0	9.4	[7.9, 10.9]	2.34	[2.12, 2.57]	No	No	No
	White	13.2	6.8	6.4	[4.5, 8.3]	1.94	[1.65, 2.24]	No	No	No
	Black	26.3	8.8	17.5	[15.0, 19.9]	2.98	[2.68, 3.30]	No	No	No
	Asian	4.4	3.4	1.0	[−0.9, 2.9]	1.30	[0.75, 1.93]	No	Yes	No
	Hispanic	30.7	8.1	22.6	[20.2, 25.1]	3.79	[3.41, 4.22]	No	No	No
	Men	12.8	5.2	7.6	[5.7, 9.5]	2.47	[2.07, 2.89]	No	No	No
	Women	19.8	8.7	11.1	[9.1, 13.1]	2.27	[2.03, 2.52]	No	No	No
	Low	81.1	13.6	67.5	[65.2, 69.9]	5.96	[5.57, 6.39]	No	No	No
	Middle	0.0	7.3	−7.3	[−9.1, −5.5]	0.00	[0.00, 0.23]	No	No	No
	High	0.0	3.7	−3.7	[−5.9, −1.6]	0.00	[0.00, 0.57]	No	No	No
GPT-4o-mini	Overall	16.9	7.0	9.9	[8.2, 11.7]	2.41	[2.15, 2.69]	No	No	No
	White	14.5	6.8	7.7	[5.3, 10.0]	2.13	[1.77, 2.50]	No	No	No
	Black	26.5	8.8	17.7	[14.8, 20.6]	3.00	[2.64, 3.38]	No	No	No
	Asian	11.2	3.4	7.8	[5.4, 10.2]	3.31	[2.50, 4.31]	No	No	No
	Hispanic	24.5	8.1	16.4	[13.7, 19.2]	3.03	[2.64, 3.45]	No	No	No
	Men	14.0	5.2	8.9	[6.6, 11.1]	2.71	[2.25, 3.20]	No	No	No
	Women	19.6	8.7	10.9	[8.3, 13.5]	2.25	[1.94, 2.57]	No	No	No
	Low	46.6	13.6	33.0	[29.8, 36.2]	3.42	[3.12, 3.76]	No	No	No
	Middle	14.4	7.3	7.1	[4.3, 9.9]	1.96	[1.57, 2.37]	No	No	No
	High	5.3	3.7	1.6	[−1.2, 4.4]	1.44	[0.69, 2.22]	No	Yes	No
GLM-4.7	Overall	13.2	7.0	6.2	[4.8, 7.7]	1.89	[1.68, 2.11]	No	No	No
	White	11.4	6.8	4.6	[2.7, 6.5]	1.68	[1.40, 1.98]	No	No	No
	Black	20.8	8.8	11.9	[9.4, 14.4]	2.35	[2.05, 2.66]	No	No	No
	Asian	3.6	3.4	0.2	[−1.7, 2.1]	1.06	[0.53, 1.66]	Yes	Yes	No
	Hispanic	20.7	8.1	12.6	[10.2, 14.9]	2.55	[2.23, 2.91]	No	No	No
	Men	7.8	5.2	2.6	[0.7, 4.5]	1.51	[1.14, 1.89]	No	Yes	No
	Women	18.3	8.7	9.6	[7.6, 11.6]	2.10	[1.86, 2.35]	No	No	No
	Low	61.4	13.6	47.8	[44.9, 50.6]	4.51	[4.18, 4.87]	No	No	No
	Middle	2.0	7.3	−5.4	[−7.4, −3.4]	0.27	[0.01, 0.53]	No	No	No
	High	0.3	3.7	−3.5	[−5.6, −1.3]	0.07	[0.00, 0.65]	No	No	No
Pooled	Overall	13.5	7.0	6.5	[5.6, 7.4]	1.92	[1.78, 2.07]	No	No	No
	White	11.3	6.8	4.5	[3.3, 5.7]	1.66	[1.48, 1.85]	No	No	No
	Black	21.9	8.8	13.1	[11.4, 14.7]	2.48	[2.27, 2.70]	No	No	No
	Asian	5.3	3.4	1.9	[0.6, 3.1]	1.55	[1.17, 2.02]	No	Yes	No
	Hispanic	21.8	8.1	13.7	[12.2, 15.2]	2.69	[2.44, 2.97]	No	No	No
	Men	9.9	5.2	4.7	[3.6, 5.8]	1.90	[1.67, 2.16]	No	No	No
	Women	16.9	8.7	8.1	[6.9, 9.4]	1.93	[1.77, 2.10]	No	No	No
	Low	55.5	13.6	41.9	[40.2, 43.5]	4.07	[3.81, 4.37]	No	No	No
	Middle	4.4	7.3	−2.9	[−4.1, −1.7]	0.61	[0.46, 0.76]	No	Yes	No
	High	1.5	3.7	−2.2	[−3.4, −0.9]	0.41	[0.10, 0.74]	No	Yes	No

Note. Post-stratified estimand, specification S1, framings combined. pp = percentage points. Ratio = simulated share over reference share, Fieller interval.

S4.4 Level 4

Figure S1

Level 2: Simulated Gaps as Ratios of the Standardised NHANES Gaps, With Intervals




Table S25*Level 3: Post-Stratified Residual of the Simulated Population Against NHANES*

Model	Overall residual, points			Groups passing 0.2 SD / 1 / 2 pt	PHQ-8 ≥ 10 (%)	
	2005–2018 [90% CI]	2021–2023	Range over 10 groups		Simulated	NHANES
DeepSeek-V3	+4.58 [4.49, 4.67]	+3.47	+4.06 to +5.13	0 / 0 / 0	7.35	7.01
Gemini-3-Flash	+2.59 [2.49, 2.69]	+1.36	+1.33 to +7.37	0 / 0 / 2	16.42	7.01
GPT-4o-mini	+4.68 [4.57, 4.79]	+3.57	+4.23 to +5.32	0 / 0 / 0	16.92	7.01
GLM-4.7	+2.12 [2.00, 2.24]	+0.91	+0.66 to +6.40	0 / 1 / 4	13.23	7.01
Pooled (descriptive)	+3.49 [3.42, 3.57]	+2.33	+2.59 to +6.05	0 / 0 / 0	13.48	7.01

Note. Residual = simulated minus NHANES mean PHQ-8, with each persona mean weighted by the NHANES share of its demographic cell (post-stratification), both framings averaged, over the 48 personas with a reference cell. The 10 groups are all adults, four races, two sexes and three income bands. Groups passing = rows whose 90% interval lies inside ± 0.79 points (0.2 SD), ± 1 point or ± 2 points; a model passes for a use only when all 10 pass. 2021–2023 = the residual against NHANES 2021–2023, analysed with its own design and with cohabiting adults counted as married (the 2005–2018 column leaves them out). The simulated share scoring 10 or more is post-stratified in the same way. Every group row is in Supplement S4.

Table S26*Level 4: Structure of the Simulated Population Against NHANES*

Model	Framing	R1: general factor			R2: loadings		R4: invariance		R3 (diagnostic)		
		λ_1/λ_2	Min. loading		ϕ_T [lower]	RMSD [upper]	at .08	Level 4	Between	Within	λ_1/λ_2
DeepSeek-V3	Clinical	1.73	-.72	Fail	.349 [.285]	.755 [.786] (Fail)	–	Fail	.17		1.45
	Narrative	1.43	-.68	Fail	.163 [.074]	.823 [.860] (Fail)	–	Fail	.14		1.34
Gemini-3-Flash	Clinical	8.69	.71	Pass	.990 [.986]	.138 [.150] Fail	Fail	Fail	.63		1.29
	Narrative	9.33	.69	Pass	.990 [.987]	.150 [.159] Fail	Fail	Fail	.64		1.45
GPT-4o-mini	Clinical	1.09	-.24	Fail	.364 [.178]	.718 [.778] (Fail)	–	Fail	.10		1.23
	Narrative	1.33	-.57	Fail	.043 [-.017]	.825 [.846] (Fail)	–	Fail	.06		1.26
GLM-4.7	Clinical	5.73	.66	Pass	.996 [.994]	.081 [.090] Pass	Fail	Fail	.48		1.59
	Narrative	8.12	.75	Pass	.998 [.997]	.100 [.112] Fail	Fail	Fail	.59		1.34
NHANES		7.47	.67						.035		4.33

Note. One-factor model on polychoric correlations of all draws of one model and framing. R1 passes when every loading is at least .30 and $\lambda_1/\lambda_2 \geq 3$. ϕ_T = Tucker’s congruence with the NHANES loadings, with the lower limit of its 90% interval; RMSD = root mean square difference from the NHANES loadings, with the upper limit of its 90% interval. R2 passes when the congruence lower limit is at least .95 and the RMSD upper limit is at most .10. R4 = every metric and scalar invariance step across sex, race and income matching the NHANES verdict at an RMSEA_D margin of .08; it fails on any determinate step that differs from NHANES, is unresolved when no step differs but some are indeterminate, and is not read for a population without a general factor (–). Parenthesised R2 verdicts are descriptive, because R1 failed. Between = share of item variance between personas (NHANES: between demographic cells). Within λ_1/λ_2 = the R1 ratio fitted to the within-persona (NHANES: within-cell) covariance. Level 4 passes on R1, R2 and R4.

Level 4 Invariance Steps: RMSEA of the Difference

Population	Attribute	Step	G	d	RMSEA $_D$	90% CI	$\hat{c}$	At .08	At .05
NHANES 2005-2018 (weighted)	sex	metric	2	7	.033	[.020, .044]	4.18	holds	holds
	sex	scalar	2	7	.067	[.061, .073]	1.67	holds	fails
	sex	metric (polychoric)	2	7	.020	[.009, .029]	3.14	holds	holds
	race	metric	4	21	.033	[.019, .043]	5.20	holds	holds
	race	scalar	4	21	.052	[.048, .056]	3.24	holds	indeterminate
	race	metric (polychoric)	4	21	.037	[.027, .047]	3.42	holds	holds
	income	metric	3	14	.051	[.037, .062]	4.58	holds	indeterminate
	income	scalar	3	14	.026	[.018, .032]	1.63	holds	holds
	income	metric (polychoric)	3	14	.031	[.021, .042]	3.16	holds	holds
DeepSeek-V3 clinical	sex	metric	2	7	.100	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	sex	scalar	2	7	.059	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	sex	metric (polychoric)	2	7	.052	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	metric	4	21	.109	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	scalar	4	21	.072	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	metric (polychoric)	4	21	.145	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	metric	3	14	.205	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	scalar	3	14	.223	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	metric (polychoric)	3	14	.243	–	–	not interpretable (no general factor)	not interpretable (no general factor)
DeepSeek-V3 narrative	sex	metric	2	7	.033	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	sex	scalar	2	7	.047	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	sex	metric (polychoric)	2	7	.330	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	metric	4	21	.139	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	scalar	4	21	.152	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	metric (polychoric)	4	21	.144	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	metric	3	14	.211	–	–	not interpretable (no general factor)	not interpretable (no general factor)

Table S27, continued

Population	Attribute	Step	G	d	RMSEA _D	90% CI	$\hat{c}$	At .08	At .05
	income	scalar	3	14	.263	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	metric (polychoric)	3	14	.252	–	–	not interpretable (no general factor)	not interpretable (no general factor)
Gemini-3-Flash clinical	sex	metric	2	7	.070	[.000, .127]	4.79	indeterminate	indeterminate
	sex	scalar	2	7	.116	[.052, .183]	4.83	indeterminate	fails
	sex	metric (polychoric)	2	7	.047	[.000, .138]	7.39	indeterminate	indeterminate
	race	metric	4	21	.000	[.000, .062]	4.39	holds	indeterminate
	race	scalar	4	21	.040	[.000, .104]	5.28	indeterminate	indeterminate
	race	metric (polychoric)	4	21	.000	[.000, .080]	5.80	holds	indeterminate
	income	metric	3	14	.192	[.151, .230]	4.19	fails	fails
	income	scalar	3	14	.096	[.061, .125]	5.49	indeterminate	fails
	income	metric (polychoric)	3	14	.148	[.053, .197]	5.74	indeterminate	fails
Gemini-3-Flash narrative	sex	metric	2	7	.039	[.000, .105]	2.85	indeterminate	indeterminate
	sex	scalar	2	7	.034	[.000, .120]	4.38	indeterminate	indeterminate
	sex	metric (polychoric)	2	7	.000	[.000, .141]	13.66	indeterminate	indeterminate
	race	metric	4	21	.000	[.000, .037]	3.88	holds	holds
	race	scalar	4	21	.000	[.000, .073]	4.23	holds	indeterminate
	race	metric (polychoric)	4	21	.069	[.000, .127]	4.84	indeterminate	indeterminate
	income	metric	3	14	.159	[.134, .186]	4.00	fails	fails
	income	scalar	3	14	.160	[.123, .196]	4.17	fails	fails
	income	metric (polychoric)	3	14	.167	[.120, .216]	5.09	fails	fails
GLM-4.7 clinical	sex	metric	2	7	.000	[.000, .072]	2.57	holds	indeterminate
	sex	scalar	2	7	.081	[.048, .108]	2.17	indeterminate	indeterminate
	sex	metric (polychoric)	2	7	.000	[.000, .059]	2.67	holds	indeterminate
	race	metric	4	21	.035	[.000, .071]	2.16	holds	indeterminate
	race	scalar	4	21	.043	[.000, .078]	2.82	holds	indeterminate
	race	metric (polychoric)	4	21	.012	[.000, .062]	2.11	holds	indeterminate
	income	metric	3	14	.162	[.139, .185]	2.28	fails	fails
	income	scalar	3	14	.160	[.136, .184]	2.87	fails	fails
	income	metric (polychoric)	3	14	.081	[.053, .105]	2.14	indeterminate	fails
GLM-4.7 narrative	sex	metric	2	7	.102	[.023, .149]	3.00	indeterminate	indeterminate
	sex	scalar	2	7	.052	[.000, .091]	2.13	indeterminate	indeterminate
	sex	metric (polychoric)	2	7	.044	[.000, .108]	3.22	indeterminate	indeterminate
	race	metric	4	21	.000	[.000, .053]	2.66	holds	indeterminate
	race	scalar	4	21	.000	[.000, .060]	3.26	holds	indeterminate
	race	metric (polychoric)	4	21	.000	[.000, .031]	2.77	holds	holds
	income	metric	3	14	.199	[.180, .219]	2.70	fails	fails
	income	scalar	3	14	.190	[.165, .212]	3.42	fails	fails
	income	metric (polychoric)	3	14	.123	[.097, .148]	2.18	fails	fails
GPT-4o-mini clinical	sex	metric	2	7	.000	–	–	not interpretable (no general factor)	not interpretable (no general factor)

Table S27, continued

Population	Attribute	Step	G	d	RMSEA _D	90% CI	$\hat{c}$	At .08	At .05
	sex	scalar	2	7	.033	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	sex	metric (polychoric)	2	7	.022	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	metric	4	21	.019	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	scalar	4	21	.035	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	metric (polychoric)	4	21	.076	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	metric	3	14	.121	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	scalar	3	14	.246	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	metric (polychoric)	3	14	.153	–	–	not interpretable (no general factor)	not interpretable (no general factor)
GPT-4o-mini narrative	sex	metric	2	7	.006	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	sex	scalar	2	7	.031	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	sex	metric (polychoric)	2	7	.028	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	metric	4	21	.048	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	scalar	4	21	.060	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	race	metric (polychoric)	4	21	.066	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	metric	3	14	.099	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	scalar	3	14	.169	–	–	not interpretable (no general factor)	not interpretable (no general factor)
	income	metric (polychoric)	3	14	.106	–	–	not interpretable (no general factor)	not interpretable (no general factor)

Note. A step holds when the upper 90% limit of RMSEA_D is below the margin, fails when the lower limit is above it, and is indeterminate otherwise. The margin of .08 is primary and .05 a sensitivity analysis. $\hat{c}$ = permutation scaling of the expected ΔF under exact invariance. The polychoric metric rows are reported as fitted.

S4.5 Multiplicity

Table S28 gives the family and the procedure at each rung. Each rung's pass is an intersection-union test over its members, except level 2, where each contrast receives its

own verdict and the family is controlled by Bonferroni.

Table S28

Multiplicity: Family and Procedure at Each Rung

Rung	Claim	Family	Members	Procedure
Gate	Draws needed per persona	Each model $\times$ framing (4×2)	Estimation of $\phi(k)$ and $SE(k)$; no test	None; persona-bootstrap intervals reported
Level 1	Person-fit like people's	Each model $\times$ framing	2 tails: misfit ratio and overfit ratio, each against $[1/\tau, \tau]$, $\tau = 1.5$	Intersection-union: pass needs both 90% intervals inside the band; no adjustment
Level 2	Group gaps like people's	Each model (the pooled row is description)	7 contrasts	Bonferroni: one-sided $\alpha = .05/7 = .0071$, 98.6% Fieller intervals
Level 3	Group levels like people's, at a stated tolerance	Each model $\times$ tolerance	10 group rows	Intersection-union: pass needs every row's 90% interval inside the tolerance (two one-sided tests at .05); no adjustment
Level 4	Structure like people's	Each model	R1, R2 and R4 in each of 2 framings; R4 reads 6 invariance steps	Intersection-union: pass needs every rule in every framing; R4 needs every step to match the reference's verdict; R3 is a diagnostic and not counted

Note. Family sizes, τ , α and the interval level are read from the receipts named in the source comment of this table.

S5. Controls

Table S29 gives, for each rung, the dose at which a planted failure is detected in at least 95% of replicates, beside the distortion the real models show.

Each replicate splits the NHANES 2005–2018 adults at random into two halves, stratified by the 48 persona cells. One half supplies respondents and the other half is the reference. A pseudo-model has the corpus design of 120 personas, two framings and 30 draws. Each draw is a respondent from the donor half, drawn within the persona's cell with probability proportional to the examination weight, and it keeps that respondent's eight item scores. Transgender and multiracial personas draw from the closest available group and enter levels 1 and 4 only. Four independent pseudo-models form a panel, and there are 200 replicates. Planted failures are applied to the same panels at a grid of doses. At level 1

Table S29*Detection of Planted Failures, Against the Distortions the Models Show*

Rung	Planted failure	Detected in $\geq 95\%$ from	Real models
Level 1	Share of draws with item values shuffled	40% of draws	Not seen
Level 1	Share of draws replaced by the most typical vector at the same total	80% of draws	Misfit ratio 0.00 to 0.18 (both framings); a 100% dose reaches 0.33
Level 2	Income or sex gap scaled by g	$g = 3$ per model (97.4% to 100%)	Low minus High, standardised: 1.9 to 5.1 per model
Level 3	Every persona mean shifted by d points	$d = 2$ at the 2-point tolerance (100%)	+2.1 to +4.7 points overall
Level 4	Share of draws decorrelated within persona	50% fails R2 in 100%; 75% fails R1 in 98.5%	Within-persona λ_1/λ_2 1.23 to 1.59, reached by the planted dose between 75% (1.71) and 100% (1.03)

Note. 200 replicates of a four-model panel of NHANES pseudo-models. Doses are nested, so a draw altered at one dose is altered at every larger dose. Per-dose detection rates for every rung and tolerance are in Supplement S5.

a share of draws has its item values shuffled, or is replaced by the most typical vector at the same total. At level 2 a group gap is scaled by a factor g . At level 3 every persona mean is shifted by d points. At level 4 a share of draws is rebuilt item by item from other draws of the same persona. Table S33 adds a parametric simulation at the corpus's own precision, with the simulated-gap standard error set to the smallest and the largest of the four real models. The design and scripts are described in `paper_brm/analysis_brm/CONTROLS.md`.

Table S30*Controls, Level 1: Detection by Dose*

Condition	Dose (%)	Unit	n	Expected reading (%)	Pass	Compressed	Excess misfit	Undet.	Misfit	Overfit
Positive control	0.0	both	800	99.4 [98.5, 99.7]	99.4	0.0	0.0	0.6	1.15	1.00
	0.0	clinical	800	97.6 [96.3, 98.5]	97.6	0.0	0.0	2.4	1.15	1.00
	0.0	narrative	800	96.4 [94.8, 97.5]	96.4	0.0	0.0	3.6	1.15	1.01
Shuffle item values	10.0	both	200	0.0 [0.0, 1.9]	28.5	0.0	0.0	71.5	1.39	0.95
	10.0	clinical	200	0.5 [0.1, 2.8]	21.5	0.0	0.5	78.0	1.39	0.95
	10.0	narrative	200	0.5 [0.1, 2.8]	24.5	0.0	0.5	75.0	1.38	0.95
	20.0	both	200	49.5 [42.6, 56.4]	0.0	0.0	49.5	50.5	1.63	0.90

Table S30, continued

Condition	Dose (%)	Unit	<i>n</i>	Expected reading (%)	Pass	Compressed	Excess misfit	Undet.	Misfit	Overfit
	20.0	clinical	200	42.5 [35.9, 49.4]	0.0	0.0	42.5	57.5	1.64	0.89
	20.0	narrative	200	35.5 [29.2, 42.3]	0.0	0.0	35.5	64.5	1.62	0.90
	40.0	both	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	2.11	0.79
	40.0	clinical	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	2.11	0.79
	40.0	narrative	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	2.11	0.79
	60.0	both	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	2.59	0.68
	60.0	clinical	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	2.59	0.68
	60.0	narrative	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	2.58	0.68
	80.0	both	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	3.07	0.58
	80.0	clinical	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	3.08	0.58
	80.0	narrative	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	3.06	0.58
	100.0	both	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	3.55	0.47
	100.0	clinical	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	3.55	0.47
	100.0	narrative	200	100.0 [98.1, 100.0]	0.0	0.0	100.0	0.0	3.54	0.47
Most typical vector	10.0	both	200	0.0 [0.0, 1.9]	100.0	0.0	0.0	0.0	1.06	0.94
	10.0	clinical	200	0.0 [0.0, 1.9]	100.0	0.0	0.0	0.0	1.07	0.94
	10.0	narrative	200	0.0 [0.0, 1.9]	100.0	0.0	0.0	0.0	1.06	0.94
	20.0	both	200	0.0 [0.0, 1.9]	100.0	0.0	0.0	0.0	0.98	0.88
	20.0	clinical	200	0.0 [0.0, 1.9]	100.0	0.0	0.0	0.0	0.98	0.88
	20.0	narrative	200	0.0 [0.0, 1.9]	99.5	0.0	0.0	0.5	0.98	0.88
	40.0	both	200	0.0 [0.0, 1.9]	72.5	0.0	0.0	27.5	0.82	0.76
	40.0	clinical	200	0.0 [0.0, 1.9]	54.0	0.0	0.0	46.0	0.82	0.76
	40.0	narrative	200	0.0 [0.0, 1.9]	52.5	0.0	0.0	47.5	0.82	0.76
	60.0	both	200	60.0 [53.1, 66.5]	0.0	2.0	0.0	40.0	0.66	0.63
	60.0	clinical	200	48.0 [41.2, 54.9]	0.0	0.5	0.0	52.0	0.65	0.63
	60.0	narrative	200	44.5 [37.8, 51.4]	0.0	1.0	0.0	55.5	0.66	0.64
	80.0	both	200	100.0 [98.1, 100.0]	0.0	80.0	0.0	0.0	0.49	0.51
	80.0	clinical	200	100.0 [98.1, 100.0]	0.0	80.0	0.0	0.0	0.49	0.51
	80.0	narrative	200	100.0 [98.1, 100.0]	0.0	80.0	0.0	0.0	0.49	0.51
	100.0	both	200	100.0 [98.1, 100.0]	0.0	89.5	0.0	0.0	0.33	0.39
	100.0	clinical	200	100.0 [98.1, 100.0]	0.0	86.5	0.0	0.0	0.33	0.39
	100.0	narrative	200	100.0 [98.1, 100.0]	0.0	88.5	0.0	0.0	0.33	0.39

Note. Dose = share of draws altered. Unit: **both** = a pseudo-model read on both framings; **clinical** and **narrative** = one framing. Expected reading = share of replicates returning the reading the planted failure should produce (a pass for the positive control), with a 95% interval. Pass, Compressed, Excess misfit and Undet. are shares of replicates in percent. Misfit and Overfit are mean ratios.

Table S31

Controls, Level 2: Verdicts by Dose

Reference, scope	Contrast	<i>g</i>	True region	<i>n</i>	Correct	Correct or compatible	Wrong	Undet.	Stopped	
half sample, per model	Black minus White	−0.50	reversed	800	0.9		2.5	0.0	0.0	97.5
		0.00	missing	800	0.0		1.0	0.0	0.0	99.0
		0.50	attenuated	800	0.0		0.0	0.4	0.0	99.6
		1.00	kept	800	0.0		0.0	0.0	0.0	100.0
		1.50	steepened	800	0.0		0.0	0.0	0.0	100.0
		2.00	steepened	800	0.1		0.1	0.0	0.0	99.9
	Hispanic minus White	3.00	steepened	800	3.6		3.6	0.0	0.0	96.4
		1.00	kept	800	0.0		0.0	0.0	0.0	100.0
	Asian minus White	1.00	kept	800	0.0		0.0	0.0	1.0	99.0
	Women minus Men	−0.50	reversed	800	4.4		75.0	0.0	3.2	21.8
		0.00	missing	800	0.0		5.0	0.0	21.0	74.0
		0.50	attenuated	800	0.0		1.4	0.0	22.2	76.4
		1.00	kept	800	0.0		0.4	0.0	22.6	77.0
		1.50	steepened	800	4.1		17.5	0.0	9.4	73.1
		2.00	steepened	800	65.5		77.9	0.0	0.1	22.0

Table S31, continued

Reference, scope	Contrast	<i>g</i>	True region	<i>n</i>	Correct	Correct or compatible	Wrong	Undet.	Stopped
	Low minus High SES	3.00	steepened	800	100.0	100.0	0.0	0.0	0.0
		−0.50	reversed	800	3.1	81.5	0.0	10.8	7.8
		0.00	missing	800	0.0	4.0	0.0	63.0	33.0
		0.50	attenuated	800	0.0	2.6	0.0	63.7	33.6
		1.00	kept	800	0.0	3.2	0.0	62.7	34.0
		1.50	steepened	800	3.4	55.4	0.0	12.8	31.9
		2.00	steepened	800	78.5	96.8	0.0	0.0	3.2
		3.00	steepened	800	100.0	100.0	0.0	0.0	0.0
	Middle minus High SES	−0.50	reversed	800	2.6	38.8	0.0	0.0	61.3
		0.00	missing	800	0.0	3.2	0.0	0.0	96.8
		0.50	attenuated	800	0.0	0.2	0.0	0.0	99.8
		1.00	kept	800	0.0	0.0	0.1	0.0	99.9
		1.50	steepened	800	3.5	3.5	0.0	0.0	96.5
		2.00	steepened	800	28.2	28.2	0.0	0.0	71.8
		3.00	steepened	800	97.4	97.4	0.0	0.0	2.6
	Low minus Middle SES	1.00	kept	800	0.0	0.0	0.1	0.0	99.9
half sample, pooled	Black minus White	−0.50	reversed	200	2.5	19.5	0.0	0.0	80.5
		0.00	missing	200	0.0	5.0	0.0	0.0	95.0
		0.50	attenuated	200	0.0	0.0	0.0	0.0	100.0
		1.00	kept	200	0.0	0.0	0.0	0.0	100.0
		1.50	steepened	200	0.0	0.0	0.0	0.0	100.0
		2.00	steepened	200	0.0	0.0	0.0	0.0	100.0
	Hispanic minus White	3.00	steepened	200	14.5	14.5	0.0	0.0	85.5
		1.00	kept	200	0.0	0.0	0.0	0.0	100.0
	Asian minus White	1.00	kept	200	0.0	0.5	0.0	0.5	99.0
	Women minus Men	−0.50	reversed	200	29.5	100.0	0.0	0.0	0.0
		0.00	missing	200	0.0	70.5	0.0	4.0	25.5
		0.50	attenuated	200	0.0	28.5	0.0	5.0	66.5
		1.00	kept	200	0.0	12.0	0.0	11.0	77.0
		1.50	steepened	200	17.0	39.5	0.0	0.0	60.5
		2.00	steepened	200	96.5	99.0	0.0	0.0	1.0
		3.00	steepened	200	100.0	100.0	0.0	0.0	0.0
	Low minus High SES	−0.50	reversed	200	41.0	100.0	0.0	0.0	0.0
		0.00	missing	200	6.0	83.0	0.0	11.0	6.0
		0.50	attenuated	200	0.0	51.0	0.0	16.0	33.0
		1.00	kept	200	0.0	32.0	0.0	34.0	34.0
		1.50	steepened	200	18.0	77.0	0.0	0.0	23.0
		2.00	steepened	200	99.5	100.0	0.0	0.0	0.0
		3.00	steepened	200	100.0	100.0	0.0	0.0	0.0
	Middle minus High SES	−0.50	reversed	200	20.5	95.5	0.0	0.0	4.5
		0.00	missing	200	0.0	29.5	0.0	0.0	70.5
		0.50	attenuated	200	0.0	12.5	0.0	0.0	87.5
		1.00	kept	200	0.0	0.0	1.0	0.0	99.0
		1.50	steepened	200	8.0	8.0	0.0	0.0	92.0
		2.00	steepened	200	58.5	58.5	0.0	0.0	41.5
		3.00	steepened	200	100.0	100.0	0.0	0.0	0.0
	Low minus Middle SES	1.00	kept	200	0.0	0.0	0.5	0.0	99.5
audit precision (approx.), per model	Black minus White	−0.50	reversed	800	0.9	3.1	0.0	0.0	96.9
		0.00	missing	800	0.0	1.0	0.0	0.0	99.0
		0.50	attenuated	800	0.0	0.0	0.4	0.0	99.6
		1.00	kept	800	0.0	0.0	0.1	0.0	99.9
		1.50	steepened	800	0.6	0.6	0.0	0.0	99.4
		2.00	steepened	800	2.4	2.4	0.0	0.0	97.6
		3.00	steepened	800	16.6	16.6	0.0	0.0	83.4

Table S31, continued

Reference, scope	Contrast	<i>g</i>	True region	<i>n</i>	Correct	Correct or compatible	Wrong	Undet.	Stopped
	Hispanic minus White	1.00	kept	800	0.0	0.0	0.1	0.0	99.9
	Asian minus White	1.00	kept	800	0.0	0.0	0.0	55.5	44.5
	Women minus Men	-0.50	reversed	800	4.6	75.4	0.0	24.4	0.2
		0.00	missing	800	0.0	5.1	0.0	94.4	0.5
		0.50	attenuated	800	0.0	2.4	0.0	97.1	0.5
		1.00	kept	800	0.0	6.5	0.0	93.0	0.5
		1.50	steepened	800	6.1	76.2	0.0	23.4	0.4
		2.00	steepened	800	70.9	99.9	0.0	0.1	0.0
		3.00	steepened	800	100.0	100.0	0.0	0.0	0.0
	Low minus High SES	-0.50	reversed	800	3.2	81.8	0.0	18.2	0.0
		0.00	missing	800	0.0	4.0	0.0	96.0	0.0
		0.50	attenuated	800	0.0	3.1	0.0	96.9	0.0
		1.00	kept	800	0.0	6.6	0.0	93.4	0.0
		1.50	steepened	800	4.9	86.1	0.0	13.9	0.0
		2.00	steepened	800	81.9	100.0	0.0	0.0	0.0
		3.00	steepened	800	100.0	100.0	0.0	0.0	0.0
	Middle minus High SES	-0.50	reversed	800	2.6	38.9	0.0	1.4	59.8
		0.00	missing	800	0.0	3.2	0.0	4.0	92.8
		0.50	attenuated	800	0.0	0.2	0.0	4.4	95.4
		1.00	kept	800	0.0	0.0	0.1	4.5	95.4
		1.50	steepened	800	5.5	5.9	0.0	4.1	90.0
		2.00	steepened	800	35.0	37.0	0.0	2.2	60.8
		3.00	steepened	800	98.1	98.4	0.0	0.0	1.6
	Low minus Middle SES	1.00	kept	800	0.0	0.0	0.2	6.0	93.8
audit precision (approx.), pooled	Black minus White	-0.50	reversed	200	2.5	23.5	0.0	0.0	76.5
		0.00	missing	200	0.0	5.0	0.0	0.0	95.0
		0.50	attenuated	200	0.0	0.0	0.0	0.0	100.0
		1.00	kept	200	0.0	0.0	0.0	0.0	100.0
		1.50	steepened	200	0.5	0.5	0.0	0.0	99.5
		2.00	steepened	200	10.0	10.0	0.0	0.0	90.0
		3.00	steepened	200	57.5	57.5	0.0	0.0	42.5
	Hispanic minus White	1.00	kept	200	0.0	0.0	0.0	0.0	100.0
	Asian minus White	1.00	kept	200	0.0	14.0	0.0	41.5	44.5
	Women minus Men	-0.50	reversed	200	32.0	100.0	0.0	0.0	0.0
		0.00	missing	200	0.0	72.5	0.0	27.5	0.0
		0.50	attenuated	200	0.0	69.5	0.0	30.0	0.5
		1.00	kept	200	0.0	57.0	0.0	42.5	0.5
		1.50	steepened	200	31.0	100.0	0.0	0.0	0.0
		2.00	steepened	200	98.5	100.0	0.0	0.0	0.0
		3.00	steepened	200	100.0	100.0	0.0	0.0	0.0
	Low minus High SES	-0.50	reversed	200	43.0	100.0	0.0	0.0	0.0
		0.00	missing	200	6.5	85.0	0.0	15.0	0.0
		0.50	attenuated	200	1.5	82.0	0.0	18.0	0.0
		1.00	kept	200	0.0	68.0	0.0	32.0	0.0
		1.50	steepened	200	34.0	100.0	0.0	0.0	0.0
		2.00	steepened	200	100.0	100.0	0.0	0.0	0.0
		3.00	steepened	200	100.0	100.0	0.0	0.0	0.0
	Middle minus High SES	-0.50	reversed	200	20.5	96.5	0.0	0.0	3.5
		0.00	missing	200	0.0	32.5	0.0	1.5	66.0
		0.50	attenuated	200	0.0	20.0	0.0	0.5	79.5
		1.00	kept	200	0.0	3.0	1.0	1.5	94.5
		1.50	steepened	200	17.5	19.0	0.0	3.0	78.0
		2.00	steepened	200	75.5	79.5	0.0	0.0	20.5
		3.00	steepened	200	100.0	100.0	0.0	0.0	0.0
	Low minus Middle SES	1.00	kept	200	0.0	1.0	0.5	5.0	93.5

Note. g = planted ratio of the simulated gap to the population gap. Shares in percent. Correct = the single true region named; Correct or compatible = a verdict that contains the true region; Wrong = a determinate verdict that excludes it; Stopped = either stop. Half sample: the reference is the other NHANES half. Audit precision: the reference standard error is scaled to the full NHANES sample.

Table S32*Controls, Level 2: Coverage of the True Ratio*

Scope	Contrast	g	n	Coverage	Mean ratio	SD of ratio
per model	Black minus White	-0.50	800	.995	-0.44	0.90
		0.00	800	.995	0.10	0.96
		0.50	800	.993	0.65	1.06
		1.00	800	.991	1.19	1.20
		1.50	800	.993	1.74	1.36
		2.00	800	.996	2.28	1.54
		3.00	800	.999	3.37	1.93
	Hispanic minus White	1.00	800	1.000	1.55	6.24
	Asian minus White	1.00	800	.998	1.02	0.28
	Women minus Men	-0.50	800	1.000	-0.49	0.17
		0.00	800	.998	0.02	0.18
		0.50	800	.995	0.52	0.20
		1.00	800	.993	1.03	0.23
		1.50	800	.993	1.53	0.25
		2.00	800	.994	2.04	0.28
		3.00	800	.995	3.05	0.35
	Low minus High SES	-0.50	800	1.000	-0.49	0.14
		0.00	800	1.000	0.01	0.15
		0.50	800	.998	0.52	0.16
		1.00	800	.998	1.02	0.18
		1.50	800	.998	1.53	0.20
		2.00	800	.998	2.03	0.23
		3.00	800	1.000	3.04	0.27
	Middle minus High SES	-0.50	800	.996	-0.49	0.28
		0.00	800	.991	0.02	0.30
		0.50	800	.988	0.53	0.33
		1.00	800	.988	1.04	0.37
		1.50	800	.986	1.55	0.42
		2.00	800	.989	2.06	0.47
		3.00	800	.991	3.08	0.57
	Low minus Middle SES	1.00	800	.988	1.05	0.38
pooled	Black minus White	-0.50	200	1.000	-0.44	0.46
		0.00	200	.990	0.10	0.57
		0.50	200	.995	0.65	0.73
		1.00	200	.995	1.19	0.92
		1.50	200	.995	1.74	1.12
		2.00	200	.995	2.28	1.33
		3.00	200	1.000	3.37	1.77
	Hispanic minus White	1.00	200	.990	1.55	5.23
	Asian minus White	1.00	200	.990	1.02	0.22
	Women minus Men	-0.50	200	1.000	-0.49	0.09
		0.00	200	.975	0.02	0.11
		0.50	200	.965	0.52	0.14

Table S32, continued

Scope	Contrast	<i>g</i>	<i>n</i>	Coverage	Mean ratio	SD of ratio
	Low minus High SES	1.00	200	.975	1.03	0.17
		1.50	200	.990	1.53	0.21
		2.00	200	.990	2.04	0.24
		3.00	200	.990	3.05	0.32
		−0.50	200	1.000	−0.49	0.08
		0.00	200	1.000	0.01	0.09
		0.50	200	.990	0.52	0.11
		1.00	200	.995	1.02	0.14
		1.50	200	.995	1.53	0.16
		2.00	200	.995	2.03	0.19
	Middle minus High SES	3.00	200	.995	3.04	0.25
		−0.50	200	1.000	−0.49	0.17
		0.00	200	.975	0.02	0.20
		0.50	200	.975	0.53	0.24
		1.00	200	.980	1.04	0.30
		1.50	200	.990	1.55	0.35
		2.00	200	.990	2.06	0.41
		3.00	200	.990	3.08	0.53
		1.00	200	.965	1.05	0.30
		Low minus Middle SES	1.00	200	.965	1.05

Note. Coverage = share of replicates whose interval contains the planted ratio; nominal .986.

Table S33

Controls, Level 2 at the Corpus's Own Precision: Parametric Power

Contrast	<i>g</i>	True region	Conditional, smallest	Conditional, largest	Pairs, smallest	Pairs, largest
Black minus White	−0.50	reversed	4.9 (63.3)	2.8 (33.7)	1.1 (23.6)	0.4 (3.1)
	0.00	missing	0.0 (8.9)	0.0 (3.8)	0.0 (1.7)	0.0 (0.5)
	0.50	attenuated	0.0 (2.0)	0.0 (0.1)	0.0 (0.0)	0.0 (0.0)
	1.00	kept	0.0 (0.0)	0.0 (0.0)	0.0 (0.0)	0.0 (0.0)
	1.50	steepened	0.1 (0.1)	0.2 (0.2)	0.1 (0.1)	0.2 (0.2)
	2.00	steepened	6.5 (6.5)	5.3 (5.3)	4.5 (4.5)	1.6 (1.6)
	3.00	steepened	72.2 (72.2)	62.6 (62.6)	61.6 (61.6)	19.5 (19.5)
Hispanic minus White	−0.50	reversed	0.5 (18.4)	0.5 (8.0)	0.2 (14.9)	0.1 (0.9)
	0.00	missing	0.0 (2.6)	0.0 (1.4)	0.0 (1.6)	0.0 (0.2)
	0.50	attenuated	0.0 (0.2)	0.0 (0.0)	0.0 (0.1)	0.0 (0.0)
	1.00	kept	0.0 (0.0)	0.0 (0.0)	0.0 (0.0)	0.0 (0.0)
	1.50	steepened	0.0 (0.0)	0.0 (0.0)	0.0 (0.0)	0.0 (0.0)
	2.00	steepened	0.0 (0.0)	0.1 (0.1)	0.0 (0.0)	0.1 (0.1)
	3.00	steepened	5.1 (5.1)	3.8 (3.8)	5.1 (5.1)	1.3 (1.3)
Asian minus White	−0.50	reversed	89.3 (100.0)	51.6 (100.0)	24.1 (100.0)	0.7 (3.9)
	0.00	missing	78.9 (100.0)	9.0 (93.4)	0.0 (48.0)	0.0 (0.5)
	0.50	attenuated	44.4 (67.8)	1.3 (58.6)	0.0 (32.2)	0.0 (0.0)
	1.00	kept	0.0 (44.3)	0.0 (29.3)	0.0 (20.2)	0.0 (0.2)
	1.50	steepened	25.9 (82.4)	19.8 (75.9)	14.1 (71.4)	0.7 (2.4)
	2.00	steepened	99.8 (100.0)	99.0 (100.0)	97.3 (100.0)	4.1 (10.7)
	3.00	steepened	100.0 (100.0)	100.0 (100.0)	100.0 (100.0)	39.7 (58.6)
Women minus Men	−0.50	reversed	98.5 (100.0)	75.9 (100.0)	81.5 (100.0)	9.6 (91.9)
	0.00	missing	97.0 (100.0)	52.0 (100.0)	63.4 (100.0)	0.0 (12.9)
	0.50	attenuated	76.8 (99.8)	31.5 (99.7)	42.1 (99.7)	0.0 (4.7)
	1.00	kept	19.9 (99.7)	0.5 (87.2)	1.9 (93.1)	0.0 (12.1)
	1.50	steepened	41.3 (100.0)	32.5 (100.0)	35.1 (100.0)	8.5 (90.1)
	2.00	steepened	100.0 (100.0)	100.0 (100.0)	100.0 (100.0)	85.2 (100.0)

Table S33, continued

Contrast	<i>g</i>	True region	Conditional, smallest	Conditional, largest	Pairs, smallest	Pairs, largest
Low minus High SES	3.00	steepened	100.0 (100.0)	100.0 (100.0)	100.0 (100.0)	100.0 (100.0)
	−0.50	reversed	99.7 (100.0)	90.1 (100.0)	62.6 (100.0)	1.7 (20.9)
	0.00	missing	99.2 (100.0)	81.0 (100.0)	25.2 (99.4)	0.0 (1.8)
	0.50	attenuated	86.0 (100.0)	61.3 (100.0)	14.1 (98.6)	0.0 (0.0)
	1.00	kept	37.4 (100.0)	12.6 (99.6)	0.2 (83.2)	0.0 (1.7)
	1.50	steepened	50.6 (100.0)	41.6 (100.0)	32.7 (100.0)	1.8 (20.3)
	2.00	steepened	100.0 (100.0)	100.0 (100.0)	100.0 (100.0)	20.4 (67.4)
	3.00	steepened	100.0 (100.0)	100.0 (100.0)	100.0 (100.0)	95.7 (99.8)
Middle minus High SES	−0.50	reversed	78.8 (100.0)	31.8 (100.0)	20.0 (99.8)	2.7 (41.6)
	0.00	missing	57.4 (99.4)	0.3 (61.2)	0.0 (37.3)	0.0 (2.8)
	0.50	attenuated	14.6 (33.4)	0.0 (19.0)	0.0 (14.8)	0.0 (0.0)
	1.00	kept	0.0 (0.5)	0.0 (0.6)	0.0 (0.7)	0.0 (0.0)
	1.50	steepened	13.4 (14.0)	10.3 (10.8)	9.4 (9.8)	3.0 (3.0)
	2.00	steepened	95.4 (96.0)	87.3 (87.8)	84.9 (85.4)	34.7 (35.1)
	3.00	steepened	100.0 (100.0)	100.0 (100.0)	100.0 (100.0)	99.6 (99.6)
	−0.50	reversed	77.1 (100.0)	33.5 (100.0)	40.3 (100.0)	0.9 (4.8)
Low minus Middle SES	0.00	missing	55.2 (99.7)	0.3 (65.7)	3.1 (78.2)	0.0 (0.8)
	0.50	attenuated	16.2 (36.9)	0.0 (21.2)	0.1 (25.9)	0.0 (0.0)
	1.00	kept	0.0 (3.3)	0.0 (2.8)	0.0 (2.7)	0.0 (0.0)
	1.50	steepened	16.0 (19.4)	11.7 (14.7)	13.4 (16.8)	0.9 (0.9)
	2.00	steepened	97.4 (99.2)	91.5 (93.8)	95.0 (97.1)	4.9 (5.2)
	3.00	steepened	100.0 (100.0)	100.0 (100.0)	100.0 (100.0)	41.3 (42.4)

Note. Each cell: percent of draws naming the true region alone, and in parentheses the percent naming a verdict that contains it. Columns: the simulated-gap standard error, set to the smallest or largest of the four real models under the draw-only (conditional) or the pairs standard error. The largest wrong-verdict rate in any cell is 0.1%.

Table S34

Controls, Level 3: Share of Replicates Passing All Ten Groups

Scope	Shift <i>d</i>	<i>n</i>	0.2 SD	1 point	2 points
per model	0.00	800	84.0 [81.3, 86.4]	98.1 [96.9, 98.9]	100.0 [99.5, 100.0]
	0.25	800	37.0 [33.7, 40.4]	82.5 [79.7, 85.0]	100.0 [99.5, 100.0]
	0.50	800	1.9 [1.1, 3.1]	27.8 [24.8, 31.0]	100.0 [99.5, 100.0]
	1.00	800	0.0 [0.0, 0.5]	0.0 [0.0, 0.5]	98.1 [96.9, 98.9]
	1.50	800	0.0 [0.0, 0.5]	0.0 [0.0, 0.5]	27.8 [24.8, 31.0]
	2.00	800	0.0 [0.0, 0.5]	0.0 [0.0, 0.5]	0.0 [0.0, 0.5]
	3.00	800	0.0 [0.0, 0.5]	0.0 [0.0, 0.5]	0.0 [0.0, 0.5]
pooled	0.00	200	100.0 [98.1, 100.0]	100.0 [98.1, 100.0]	100.0 [98.1, 100.0]
	0.25	200	91.5 [86.8, 94.6]	100.0 [98.1, 100.0]	100.0 [98.1, 100.0]
	0.50	200	14.0 [9.9, 19.5]	84.5 [78.8, 88.9]	100.0 [98.1, 100.0]
	1.00	200	0.0 [0.0, 1.9]	0.0 [0.0, 1.9]	100.0 [98.1, 100.0]
	1.50	200	0.0 [0.0, 1.9]	0.0 [0.0, 1.9]	84.5 [78.8, 88.9]
	2.00	200	0.0 [0.0, 1.9]	0.0 [0.0, 1.9]	0.0 [0.0, 1.9]
	3.00	200	0.0 [0.0, 1.9]	0.0 [0.0, 1.9]	0.0 [0.0, 1.9]

Note. Shift *d* = points added to every persona mean. Shares in percent with 95% intervals. At *d* = 0 the share is the pass rate of real people; at *d* > 0 a failure is the correct reading.

Table S35*Controls, Level 4: Share of Replicates Passing Each Rule*

Condition	Dose (%)	Unit	200 replicates			R2 calibration							λ_1/λ_2	RMSD	Within λ_1/λ_2
			<i>n</i>	R1	R3	<i>n</i>	R2	pass	R2 unres.	R2 fail	Pass				
L4 decorrelate	25.0	clinical	200	100.0	100.0	40	55.0		45.0	0.0	–	5.68	0.08		3.46
L4 decorrelate	25.0	model	200	100.0	100.0	40	37.5		62.5	0.0	37.5	–	–		–
L4 decorrelate	25.0	narrative	200	100.0	100.0	40	60.0		40.0	0.0	–	5.66	0.08		3.46
L4 decorrelate	50.0	clinical	200	100.0	0.0	40	0.0		0.0	100.0	–	4.13	0.17		2.55
L4 decorrelate	50.0	model	200	100.0	0.0	40	0.0		0.0	100.0	0.0	–	–		–
L4 decorrelate	50.0	narrative	200	100.0	0.0	40	0.0		0.0	100.0	–	4.13	0.17		2.56
L4 decorrelate	75.0	clinical	200	7.5	0.0	40	0.0		0.0	100.0	–	2.75	0.29		1.71
L4 decorrelate	75.0	model	200	1.5	0.0	40	0.0		0.0	100.0	0.0	–	–		–
L4 decorrelate	75.0	narrative	200	8.5	0.0	40	0.0		0.0	100.0	–	2.78	0.29		1.72
L4 decorrelate	100.0	clinical	200	0.0	0.0	40	0.0		0.0	100.0	–	1.50	0.48		1.03
L4 decorrelate	100.0	model	200	0.0	0.0	40	0.0		0.0	100.0	0.0	–	–		–
L4 decorrelate	100.0	narrative	200	0.0	0.0	40	0.0		0.0	100.0	–	1.51	0.48		1.03
Positive control	0.0	clinical	200	100.0	100.0	40	100.0		0.0	0.0	–	7.56	0.03		4.44
Positive control	0.0	model	200	100.0	100.0	40	100.0		0.0	0.0	100.0	–	–		–
Positive control	0.0	narrative	200	100.0	100.0	40	100.0		0.0	0.0	–	7.53	0.03		4.45
reference half	0.0	reference	200	100.0	100.0	–	–		–	–	–	–	–		4.25

Note. Dose = share of draws rebuilt item by item from other draws of the same persona. Rule columns are shares of replicates in percent. R2 is read on congruence and loading RMSD (pass, unresolved, fail) in the R2 calibration replicates (script 83h), which rerun the level-4 block with the RMSD recorded. Pass is scored on R1, R2 and R3 in both framings (unit **model**) in those replicates. R4 was not simulated. λ_1/λ_2 , RMSD (mean loading RMSD against the reference) and the within-persona ratio are means over replicates. The reference-half row reads the NHANES half used as reference.

S6. External Level 3: An Illustration

This section is an illustration of level 3 on two external releases, and it is not a verdict on either. Level 3 needs a use tolerance that the deployer states, and neither group of original authors stated one. The use tolerances below are placeholders set by the present analyst. On OpinionQA a group passes when the model’s plurality answer matches the Pew plurality answer on at least 90% of questions. On Argyle et al. Study 3 the use tolerance is 5 percentage points for proportions and one response category for ordinal scales. The strict tolerance is the human group mean’s own 95% sampling half-width.

OpinionQA. The steered model distributions and the Pew American Trends Panel targets come from the Meister et al. release. The Pew targets in that release are unweighted counts, so the comparison is to the unweighted Pew sample. Of 268 method, model, steering and group cells, 0 pass the placeholder use tolerance (Table S36). The models match the Pew plurality answer on 25.5% to 79.7% of questions (median 57.1%), and a fresh Pew sample of the same size would match it on a median of 94.5%.

Argyle et al. Study 3. GPT-3 answered for the same American National Election Studies respondents whose answers form the reference, and the ANES file carries no weights. The released listing for Study 3 sets `temperature=0.001` inside the generation loop, while the release’s README calls the main file the 0.7 run. This analysis follows the README and labels that run `t0.7_main`. In the main run, 4 of 116 outcome and group cells pass the strict tolerance and 57 pass the placeholder use tolerance (Tables S37 and S38). The receipts are `paper_brm/external/results/opinionqa/level3_groups.csv` and `paper_brm/external/results/argyle/level3_groups.csv`, and the method is described in `paper_brm/external/RESULTS.md`.

Table S36*External Level 3 Illustration: OpinionQA Steered Distributions Against Pew*

Method	Model	Steering	Groups	Plurality match, median (range)	Human resample	Within strict	Mean R	Groups passing
express_ distribution	anthropic_haiku	few-shot	14	57.1 (47.9–66.3)	94.3	16.8	0.10	0
express_ distribution	anthropic_haiku	persona	14	60.2 (36.6–62.7)	94.3	13.8	0.12	0
express_ distribution	anthropic_opus	few-shot	14	66.3 (61.8–78.6)	94.3	18.5	0.08	0
express_ distribution	anthropic_opus	persona	14	67.0 (45.6–79.7)	94.5	20.3	0.09	0
express_ distribution	gpt-3.5-turbo-0125	few-shot	14	46.2 (39.8–52.0)	94.3	10.3	0.16	0
express_ distribution	gpt-3.5-turbo-0125	persona	14	46.6 (26.5–51.6)	94.7	12.8	0.13	0
express_ distribution	gpt-4	few-shot	14	60.5 (54.1–68.4)	94.3	18.9	0.09	0
express_ distribution	gpt-4	persona	14	63.0 (52.8–74.5)	94.3	15.3	0.09	0
express_ distribution	llama3-70b	few-shot	6	58.2 (56.1–66.3)	96.2	14.3	0.11	0
express_ distribution	llama3-70b	persona	6	57.7 (50.0–63.3)	96.2	11.7	0.11	0
model_logprobs	gpt-3.5-turbo-0125	few-shot	14	46.0 (40.2–50.0)	94.3	5.1	0.23	0
model_logprobs	gpt-3.5-turbo-0125	persona	14	45.4 (39.8–54.1)	94.3	12.3	0.16	0
model_logprobs	gpt-4	few-shot	14	66.0 (57.7–73.5)	94.3	6.7	0.20	0
model_logprobs	gpt-4	persona	14	60.5 (43.7–69.4)	94.3	4.6	0.24	0
model_logprobs	llama3-70b	few-shot	6	60.2 (54.6–63.3)	95.3	6.1	0.19	0
model_logprobs	llama3-70b	persona	6	62.2 (52.0–69.4)	95.3	3.6	0.22	0
sequence	anthropic_haiku	few-shot	8	54.1 (47.4–59.2)	94.3	11.3	0.12	0
sequence	anthropic_haiku	persona	8	49.0 (36.7–55.1)	94.3	11.2	0.13	0
sequence	anthropic_opus	few-shot	8	64.8 (58.2–74.5)	94.3	10.7	0.12	0
sequence	anthropic_opus	persona	8	58.2 (48.0–66.3)	94.3	11.3	0.13	0
sequence	gpt-3.5-turbo-0125	few-shot	8	45.4 (41.8–50.0)	94.3	11.3	0.16	0

Table S36, continued

Method	Model	Steering	Groups	Plurality match, median (range)	Human resample	Within strict	Mean R	Groups passing
sequence	gpt-3.5-turbo-0125	persona	8	34.2 (25.5–40.8)	94.3	6.7	0.16	0
sequence	gpt-4	few-shot	8	60.2 (56.7–68.4)	94.3	15.8	0.10	0
sequence	gpt-4	persona	8	55.1 (49.5–60.2)	94.3	13.8	0.11	0
sequence	llama3-70b	few-shot	6	58.2 (51.0–66.3)	95.3	9.7	0.16	0
sequence	llama3-70b	persona	6	58.7 (56.7–63.3)	95.3	7.1	0.16	0

Note. Shares in percent, medians over groups. Plurality match = share of questions on which the model's plurality answer equals the Pew plurality answer. Human resample = the same share for a fresh Pew sample of the same size. Within strict = share of questions whose residual lies inside the human group's own 95% sampling half-width. Mean |R| = mean absolute residual as a share of the scale range. Groups passing = groups whose plurality match is at least 90%, the placeholder use tolerance.

Table S37

External Level 3 Illustration: Argyle et al. Study 3, Passing Cells by Run

Run	Outcome × group cells	Pass strict	Pass use
t0.001	116	8	51
t0.7_main	116	4	57
t1.0	116	1	58

Note. Strict tolerance: the human group mean's own 95% sampling half-width. Use tolerance: a placeholder of 5 percentage points for proportions and one response category for ordinal scales.

Table S38

External Level 3 Illustration: Argyle et al. Study 3, Main Run by Outcome

Outcome	All respondents					Groups				
	Human	Silicon	Residual	90% CI	Use tol.	Pass	n	Strict	Use	Agree (%)
Trump share among voters	0.46	0.26	−0.20	[−0.22, −0.18]	0.05	No	13	0	2	72.3
Trump two-party share	0.52	0.50	−0.02	[−0.04, −0.01]	0.05	Yes	12	3	4	86.2
Turnout 2016	0.87	0.83	−0.04	[−0.05, −0.03]	0.05	No	13	1	1	80.3
Attends church	0.60	0.60	0.00	[−0.02, 0.02]	0.05	Yes	13	0	1	52.4
Ever discusses politics	0.82	0.88	0.06	[0.05, 0.08]	0.05	No	13	0	1	76.9
Party ID (1-7)	3.85	4.39	0.54	[0.49, 0.60]	1.00	Yes	13	0	11	23.6
Ideology (1-7)	4.16	4.04	−0.12	[−0.17, −0.07]	1.00	Yes	13	0	13	22.0
Patriotism (1-7, 1 = flag feels extremely good)	1.95	1.43	−0.52	[−0.57, −0.48]	1.00	Yes	13	0	11	44.6
Political interest (1-4, 1 = very)	2.14	1.74	−0.40	[−0.43, −0.37]	1.00	Yes	13	0	13	32.6

Note. Main run (t0.7_main). Proportions for the first five outcomes, scale points for the rest. Groups: all respondents plus twelve age, education, gender, party and race groups. Agree = share of respondents whose silicon answer equals their own answer.

S7. Ladder Specification

This section gives every estimator, interval and pass rule of the ladder in full. The article states each rule in short form, and LADDER_SPEC.md in the repository is the frozen specification with the code that implements it. At level 2 the article reports three verdicts

per contrast. A contrast is KEPT when its interval lies in the kept region below, NOT KEPT when its interval excludes that region, and UNRESOLVED when the interval overlaps it. The five regions below name the direction of a contrast that is not kept, and they are the verdict labels in the receipt files.

Gate: Regeneration Stability

Question. How many draws must be averaged before a persona's score is precise enough to read?

Statistic. The gate is a generalizability study (??). Within one model and one framing, the score of draw r of persona p is

$$X_{pr} = \mu + \nu_p + e_{r:p}, \quad (1)$$

with persona variance σ_p^2 and within-cell draw variance σ_r^2 . The dependability of a mean of k draws is

$$\phi(k) = \frac{\sigma_p^2}{\sigma_p^2 + \sigma_r^2/k}, \quad (2)$$

and its standard error is $SE(k) = \sqrt{\sigma_r^2/k}$. The k needed for a target ϕ^* is

$\lceil \phi^*/(1 - \phi^*) \cdot \sigma_r^2/\sigma_p^2 \rceil$, and the k needed for a tolerance t on the standard error is $\lceil \sigma_r^2/t^2 \rceil$.

Draws of one persona carry no order, and a crossed fit of persona by draw index finds no draw main effect, so the absolute and relative coefficients coincide. Components are estimated from expected mean squares and checked against restricted maximum likelihood. Intervals are percentile bootstrap intervals over personas ($B = 1,000$). When a persona's score must also hold across prompt wording, framing becomes a random facet, and $\phi_F(k) = \sigma_p^2/(\sigma_p^2 + \sigma_f^2 + \sigma_{pf}^2 + \sigma_r^2/k)$ has a ceiling that no number of draws can pass.

Pass rule. The gate reads precision. Persona means of k draws may be read when $SE(k)$ is at most 0.25 reference standard deviations in each framing (the minimum rule), and the recommended k brings it to 0.125 standard deviations. A single draw may be read as the persona's score when the within-cell standard deviation is at most 0.25 reference

standard deviations. On the PHQ-8 these tolerances are 0.98 and 0.49 points. One point is about the size of the population's sex gap, and half a point keeps a persona mean's error below both that gap and half the low-to-high income gap. When a later level reads the persona mean over both framings, the gate is also reported for that mean. The dependability $\phi(k)$ measures how well persona means separate personas from one another. It is reported beside every verdict, and it is required ($\phi(k) \geq .80$, or .90 for the recommended k) only for uses that compare or rank individual personas. It is not a pass condition for levels 2 and 3, because those levels carry the draw error in their own intervals and because a faithful simulator has the population's small differences between cells: real NHANES respondents drawn through the persona design reach $\phi(30)$ of .63 to .76 across replicates (mean .70). The probability that two draws of one persona fall in different severity categories (band flip) is reported beside the pass rule as a description of single-draw labels.

What a pass supports. A pass at k supports reading the persona mean of k draws as the model's answer for that persona, to within the stated error. A pass with $\phi(k) \geq .80$ also supports comparing individual personas with each other. Levels 2 and 3 compare these persona means with the population. The gate says nothing about whether that answer is correct.

Level 1: Individual Coherence

Question. Is a single draw an answer pattern that a real respondent with the same total score could plausibly give, and are the draws as varied as real respondents?

Statistic. A graded response model (?) is fitted to the population reference with survey weights, and every respondent and every draw is scored on its item parameters. For an answer vector $\mathbf{x}$ with trait estimate $\hat{\theta}$, the log-likelihood person-fit statistic (?) is

$$l_0 = \sum_j \sum_k d_{jk} \log P_{jk}(\hat{\theta}), \quad d_{jk} = 1 \text{ if } x_j = k, \quad (3)$$

where P_{jk} is the model's probability of category k on item j . We standardise it with the correction for an estimated trait (?) in its polytomous form (?), giving l_z^* , with $\hat{\theta}$ the weighted likelihood estimate (?). Low l_z^* marks a vector less likely than the model expects at its trait level (misfit), and high l_z^* marks a vector more regular than expected (overfit). Because l_z^* depends on the total score, each draw is placed in the weighted reference distribution of l_z^* among respondents with the same total, using a randomised probability integral transform. The reference therefore puts exactly 5% of its own respondents below its within-total 5th percentile and 5% above its 95th, whatever the mix of totals. The *misfit ratio* is a model's share of draws below the within-total 5th percentile divided by .05, and the *overfit ratio* is its share above the 95th percentile divided by .05. A real population has both ratios equal to 1.

Pass rule. A model passes in a framing when the 90% intervals of both ratios lie inside $[1/\tau, \tau]$. The tolerance τ is the smallest value in $\{1.50, 2.00, 3.00\}$ that contains the departure of every reference subgroup (by sex, race and ethnicity, income band and survey cycle) when each subgroup is scored against the pooled reference. A subgroup's departure is read on the 90% interval of each ratio: it is the interval end nearest 1 when the interval excludes 1, and there is none when the interval covers 1. Reading the interval and not the point ratio keeps the sampling noise of a small subgroup from setting τ . The floor of 1.5 comes from the positive control in the article's Controls section: real respondents drawn through the persona design pass at 1.5 and not at 1.25, so a reference too small to resolve its subgroups does not set a tighter tolerance than real people meet. On NHANES the largest departure is 1.27, the lower end of the misfit ratio for non-Hispanic Black respondents, so $\tau = 1.5$, and every subgroup passes at that tolerance. Each tail is read on its interval: above the band, below it, inside it, or crossing a boundary. A deficit in both tails, too few atypical and too few highly regular vectors, is called *compressed*. A pass needs both tails, so the rule is an intersection-union test and needs no multiplicity adjustment (?). Intervals come from a persona-clustered bootstrap of the simulated draws,

drawn jointly with a Rao–Wu bootstrap over the reference’s primary sampling units (??) ($B = 2,000$). For the PHQ-8, the gateway rule (an elevated total with neither of the two core items at 2 or more) is reported at matched totals as a descriptive row. The row is descriptive because, with the corpus’s totals kept, uniformly random answer vectors read as too prototypical on it (observed to expected 0.62), while person fit reads them as excess misfit (ratio 6.05).

What a pass supports. A pass supports using single draws as individual cases whose item pattern is as plausible, and as varied, as a real respondent’s at the same severity. Examples are exemplar vignettes, item-level test cases for a screening tool, and training data for a classifier that reads item patterns. A compressed failure rules out uses that depend on the spread of presentations, such as stress-testing a screener with atypical cases.

Level 2: Subgroup Fidelity

Question. Is the gap between two groups in the simulated sample the same size and direction as the gap in the population?

Estimand. A simulated contrast compares persona pairs that differ on one attribute and match on the rest, so it is a conditional gap. For a contrast between groups A and B over strata s of the other attributes,

$$g = \frac{1}{S} \sum_{s=1}^S (\bar{X}_{As} - \bar{X}_{Bs}), \quad (4)$$

where $\bar{X}$ is a persona mean. The *standardised* population gap γ is the reference contrast computed within the same strata and averaged with the same equal weights. This gap answers the question the persona design asks, and it is the primary reference. The *marginal* population gap, the crude difference between the two groups in the survey, is reported beside it. The two can differ in sign when the groups differ in composition.

Statistic. The statistic is the ratio $\rho = g/\gamma$. Its interval is the Fieller set (?): the values c at which the one-sided tests of $\rho > c$ and $\rho < c$, using

$$t_c = \frac{g - c\gamma}{\sqrt{\text{SE}_g^2 + c^2 \text{SE}_\gamma^2}}, \quad (5)$$

do not reject. The simulation and the survey are independent samples, and degrees of freedom are Welch–Satterthwaite, combining the simulation’s and the survey design’s. The standard error of g is the standard deviation of the persona-pair differences divided by the square root of the number of pairs. This standard error treats differences in the gap between strata as noise, so it is conservative. A draw-only standard error, which holds the persona set fixed, is reported as a sensitivity analysis. The population gap’s standard error is a design-based Taylor linearisation of the whole contrast.

Pass rule. The boundaries -0.25 , 0.25 , 0.75 and 1.25 divide the ratio into five regions, named REVERSED, MISSING, ATTENUATED, KEPT and STEEPENED. The verdict names the region that holds the interval, names two adjacent regions when the interval crosses one boundary, and is UNDETERMINED when it crosses more. A contrast passes when it is KEPT. A model passes level 2 when every contrast that is not stopped is KEPT, fails when any verdict excludes KEPT, and is otherwise unresolved. Each model is one family of contrasts, and every one-sided test runs at .05 divided by the family size (Bonferroni), so with seven contrasts the intervals are 98.6% intervals. Two stops are checked first. The rung records NO POPULATION GAP when the population gap’s own interval covers zero, since a ratio to an unresolved denominator has no meaning. It records REFERENCE TOO IMPRECISE when the population gap’s interval is too wide for any simulated gap to be read as KEPT. With k the relative half-width of that interval, $k = t \text{SE}_\gamma / |\gamma|$, even an exactly known simulated gap gives a ratio interval spanning a factor of $(1 + k)/(1 - k)$, and the KEPT region spans a factor of $1.25/0.75$, so KEPT is out of reach when $k > 1/4$. A contrast with $k > 1/4$ whose verdict already excludes KEPT keeps that verdict, so an imprecise

reference can still show a failure but cannot hold a model at unresolved.

What a pass supports. A pass supports reading the size and direction of the group difference from the simulated sample. A kept gap says nothing about either group's level.

Level 3: Population Calibration

Question. Reweighted to the real composition of a group, does the simulated sample have the group's level?

Statistic. Each persona cell c in a group G receives the weight $p_c = N_c / \sum_{c' \in G} N_{c'}$, where N_c is the weighted count of reference respondents in that cell. The residual is

$$R_G = \sum_{c \in G} p_c (s_c - y_c), \quad (6)$$

where s_c is the persona mean and y_c is the reference mean in the same cell, so the simulated level is post-stratified to the population's composition (?). The persona cells are fixed by design, so the simulated side's sampling error is the draw sampling within cells, $\text{var}_{\text{sim}}(R_G) = \sum_c p_c^2 \text{var}(s_c)$. The reference side's error, which includes the uncertainty in the weights p_c , comes from a stratified delete-one-PSU jackknife over the survey design (?). The two sides are independent, and degrees of freedom follow Satterthwaite with the design degrees of freedom on the reference term (?). The share of the group scoring at or above the instrument's cut point is a second row, compared both as a difference in percentage points and as a ratio with a Fieller interval.

Pass rule. A row passes at tolerance δ when two one-sided tests at $\alpha = .05$ reject both $-\delta$ and $+\delta$, that is, when the 90% interval of R_G lies inside $(-\delta, +\delta)$ (??). A model passes for a use when every group row passes, and this rule is again an intersection-union test. The tolerance is the smallest difference that matters for the stated use, and the researcher names it. For the PHQ-8 we report 0.2 reference standard deviations, 1 point and 2 points, with 2 points (about half a standard deviation) as the named default. The

prevalence row uses differences of 2.5 and 5 percentage points and a ratio band of 0.80 to 1.25. The 5-point minimal important change of the PHQ-9 (?) is a within-person threshold and is not used as a population tolerance.

What a pass supports. A pass at δ supports using the simulated level as a stand-in for the group's level, within δ , for this prompt, this persona grid and this survey period. A level-3 pass on the mean does not support prevalence or screening uses, and the prevalence row tests those separately.

Level 4: Structural Fidelity

Question. Do the items hang together as they do in people, and in the same way across groups?

Unit. A model's simulated population in one framing is every draw across its personas. Its item correlations, like the population's, mix variation between demographic profiles with variation among respondents who share a profile. The persona is the resampling unit for every interval, because its draws share one prompt.

Statistics and pass rule. Polychoric correlations are estimated by two-step maximum likelihood (?), and a one-factor model is fitted by unweighted least squares. Level 4 has three scored rules and one diagnostic.

1. *R1, a general factor.* Every loading is at least .30 and the ratio of the first to the second eigenvalue, λ_1/λ_2 , is at least 3. Where the reference itself falls below a threshold, as a weak item of an established scale can, the threshold for that item or ratio is the reference's own value, so a model needs a general factor at least as strong as people's. A population without a general factor fails level 4, and its later rules are reported as description.
2. *R2, loadings like people's.* Tucker's congruence between the model's loadings λ and the reference loadings μ ,

$$\phi_T = \frac{\sum_i \lambda_i \mu_i}{\sqrt{\sum_i \lambda_i^2 \sum_i \mu_i^2}}, \quad (7)$$

has a lower 90% limit of at least .95 (?), and the root mean square difference between λ and μ has an upper 90% limit of at most .10. Congruence reads the pattern of the loadings and ignores their size, so a population whose items correlate half as strongly as people's, or more strongly, can still reach a congruence near 1; the second condition reads the size. A difference of .10 in every loading of an eight-item scale with loadings of .70 moves coefficient omega from .885 to .818 (loadings of .60) or .934 (loadings of .80). Both limits come from the persona bootstrap of the model paired with the bootstrap of the reference.

3. *R4, invariance like people's.* Multi-group confirmatory factor models test metric and scalar invariance across sex, race and income. The change in fit between nested steps is read as the root mean square error of approximation of the difference (?),

$$\text{RMSEA}_D = \sqrt{G} \sqrt{\max(0, \Delta F - \nu)/d}, \quad (8)$$

where ΔF is the change in the fit function, d the change in degrees of freedom, G the number of groups and ν the expected ΔF under exact invariance. Draws are clustered in personas and the items are coarse, so ν is estimated by permuting group labels among personas (?). A step holds when the upper limit of its 90% interval is below the margin, fails when the lower limit is above it, and is indeterminate otherwise (??). The margin is .08, because the reference itself fails or is indeterminate on three of its six steps at .05; .05 is a sensitivity analysis. R4 passes when every model step matches the reference's verdict on the same step, fails when any determinate step differs from it, and is unresolved when no step differs but some are indeterminate. A model with an indeterminate step therefore cannot pass R4, and a design needs enough personas per group to resolve every step.

4. *R3, diagnostic: the source of the factor.* Each population's item covariance is split into a between-cell part and a pooled within-cell part, where a cell is a persona in the

simulation and a demographic profile in the reference. The one-factor model is fitted to the within-cell part. This shows whether the factor lives among draws of one persona, as it lives among people who share a profile, and it feeds the same restriction as level 1.

A model passes level 4 when it passes R1, R2 and R4 in every framing.

What a pass supports. Passing R1 and R2 supports scoring the simulated population's total as one score with human-like item weights and human-like reliability. Passing R4 adds that a group comparison of the total means the same thing in the simulation as in people. A within-persona factor (R3) would support treating draws as individual respondents in item-level psychometric work.